

Procedural and learning-based generation of coral reef islands

Submission1265

Abstract

We propose a procedural method for generating single volcanic islands with coral reefs using user sketching from two projections: a top view, which defines the island's shape, and a profile view, which outlines its elevation. These projections, commonly used in geological and remote sensing domains, are complemented by a user-defined wind field, applied as a distortion field to deform the island's shape, mimicking the effects of wind and waves on the long term and enabling finer user control. We then model the growth of coral on the island and its surrounding to construct the reef following biological observations. Based on these inputs, our method generates a height field of the island. Our method is capable of creating a large variety of island models composing a dataset used for training a conditional Generative Adversarial Network (cGAN). By applying data augmentation, the cGAN allows for even greater variety in the generated islands, providing users with higher freedom and intuitive controls over the shape and structure of the final output.

Keywords: Procedural modeling, Terrain synthesis, cGAN, Coral reef, Sketch-based interface

1. Introduction

Simulating the formation of coral reef islands presents significant challenges due to the complex interplay of geological, environmental, and biological factors Hopley, 2014. One major difficulty lies in capturing the long-term subsidence of volcanic islands, which occurs over millions of years, while simultaneously modeling the upward growth of coral reefs that rely on environmental conditions such as water depth, temperature, and sunlight. This combination of slow geological processes and dynamic biological growth is difficult to replicate in a computational model.

Additionally, the biological aspects of coral growth are inherently tied to environmental factors. Coral reefs grow only within a specific range of water depth and sunlight, and their growth patterns are affected by the health of the reef ecosystem and the availability of resources. Accurately modeling these biological dependencies in a procedural system is challenging, as these factors are numerous and difficult to generalize. Moreover, the scarcity of data available obstructs the global understanding of these biomes. In a recent high-resolution mapping of shallow coral reefs Lyons et al., 2024, researchers estimated the total surface area of this biome to cover less than 0.7% of Earth's area, and more specifically that coral habitat represents less than 0.2%.

Existing terrain generation methods, such as Perlin noise-based algorithms or uplift-erosion models, are often ill-suited for these processes. While they can generate natural-looking landscapes (such as alpine landscape, representing about a quarter of land area Körner et al., 2014), they do not account for the unique geological and biological interactions that govern coral reef island formation,

thus missing coherency. Capturing these dynamics, while also providing user control during the modeling of a terrain, requires a balance between realism and procedural flexibility, allowing for both accurate computationally expensive simulation of natural processes and intuitive user control in interactive time.

Despite advances in terrain generation, existing methods struggle with user-controlled design of specific island shapes and achieving realism without real data. Coral reef islands exemplify this gap: we lack datasets to directly train deep models, and purely procedural methods require expert tuning to mimic their features.

To address these issues, we use procedural generation as an initial step in our approach as a means to efficiently create a large and diverse set of training examples for a learning-based model. Each synthetic example is represented by a terrain height field and a corresponding semantic label map that marks different regions, providing structured input-output pairs for the learning stage as presented in ??.

We trained a conditional Generative Adversarial Network (cGAN) as the core of our learning-based approach. A cGAN is a type of deep learning model that learns to generate realistic data based on an input condition or context. In our case, the cGAN takes as input the semantic label map of an island generated by the procedural step and learns to produce a realistic island height field that matches this layout. By training on the many examples from the procedural generator, the cGAN captures the subtle terrain features and variations characteristic of coral reef islands, going beyond what hard-coded procedural rules can achieve thanks to the application of data augmentation. The cGAN model can be used on its own to generate new island terrains with simplified and more intuitive user inputs through digital drawing, and the model will generate a realistic island terrain accordingly.

The key contributions are as follows: 1) a novel sketch-based procedural algorithm for shaping island terrains from top and profile views, 2) The training of a deep learning model on synthetic data derived from procedural rules, serving as an abstraction layer that hides underlying complexity, 3) A demonstration that the cGAN approach tolerates imprecise, low-detail user input sketches, broadening usability, without the need for cutting-edge network architectures, and 4) An insight that procedural generation remains essential to produce training data in data-sparse domains such as coral reef islands. These contributions collectively show a pathway to blend user-driven design with learning-based generation in terrain modeling.

2. Related Work

Procedural terrain generation spans a spectrum from purely noise-driven algorithms to physics-based simulations and, more recently, data-driven methods. In the context of coral reef islands, where both long-term geological subsidence and biogenic reef accretion interplay, existing approaches fall short in one of two ways: either they lack ecological or geological grounding, or they offer insufficient authoring control.

Procedural and simulation-based methods Noise-based techniques such as Perlin noise Perlin, 1985, Simplex noise Perlin, 2001, and the Diamond-Square algorithm Fournier et al., 1982 (often extended via fBm or multifractal noise Musgrave et al., 1989; Ebert et al., 2003) remain popular for their simplicity and speed. Island shapes are generated by modulating noise with radial falloff masks Olsen, 2004, but these methods cannot reproduce reef rings, lagoons, or atoll structures in a geologically coherent manner. They treat terrain purely as a signal-processing problem, divorced from processes like volcanic subsidence or coral growth Smelik et al., 2009; Galin et al., 2019.

Simulation-based approaches introduce causality by modeling erosion Beneš et al., 2006; Neidhold et al., 2005; Mei et al., 2007, tectonic uplift Cordonnier, Braun, et al., 2016; Cordonnier, Cani, et al., 2017; Schott et al., 2023, or vegetation-terrain feedback Ecormier-Nocca et al., 2021; Cordonnier, Galin, et al., 2017. Hydraulic and thermal erosion capture fluvial networks and slope-driven mass wasting, but they omit underwater sedimentation and biogenic carbonate accretion. Tectonic and isostatic models excel at orogeny but ignore coral reef dynamics, while vegetation-based methods do not generalize to marine ecosystems. Consequently, none of these simulations jointly capture the slow subsidence of volcanic islands and the compensatory growth of surrounding reefs on the timescales required for atoll formation.

Sketch-based terrain modeling Sketch-driven interfaces bridge user intent and procedural detail. Curve-based systems let users draw ridges, valleys, or coastlines that guide multiresolution deformation and noise propagation Gain et al., 2009; Hnaidi et al., 2010. Constraint-based methods extend this by enforcing absolute elevation or slope values at control curves or points, solved via diffusion or fractal interpolation Gasch et al., 2020; Talgorn and Belhadj, 2018, and even gradient-domain editing for seamless slope control Guérin, Peytavie, et al., 2022. Semantic approaches encode high-level “terrain atoms” from a dictionary of primitives Génevaux et

al., 2015 or interpret geological schematics into 3D models Natali et al., 2012. While these techniques grant artists fine-grained control, they typically lack ecological constraints and have not been tailored to marine biogeomorphology.

Learning-based terrain synthesis Deep generative models offer a way to learn complex patterns without manual procedural rules. Unconditional GANs have been applied to digital elevation maps of mountains Wulff-Jensen et al., 2018 and joint height-texture synthesis Spick and Walker, 2019, but their reliance on latent noise prevents precise layout control. Two-stage pipelines use an initial GAN for heightmaps and a conditional refinement for textures Beckham and Pal, 2017 or reverse the order (imagery-to-DEM) Panagiotou and Charou, 2020, yet still lack semantic inputs.

Conditional GANs (cGANs) extend image-to-image translation methods such as pix2pix Isola et al., 2017 to terrain, enabling sketch- or label-map-conditioned generation. Prior work includes sketch-to-DEM translation for generic landforms Guérin, Digne, Galin, Peytavie, et al., 2017 and sparse “altitude dot” conditioning Voulgaris et al., 2021, as well as cGANs that invert satellite imagery into elevation Sisodia, 2022. However, these models require extensive paired datasets which are scarce for coral reef islands.

3. Overview of our method

Our method for generating coral reef islands combines user-driven sketching, procedural techniques, and deep learning to create realistic and varied island terrains (Figure 1).

The pipeline consists of two distinct phases: a procedural data-generation phase and a deep-learning-driven inference phase.

3.1. Procedural generation phase

In the initial procedural phase, the user sketches key island features from two complementary viewpoints: a top view, defining the horizontal layout of island features (island boundaries, beach width, lagoon areas, coral reefs), and a profile view, specifying the vertical elevation profile from island center to ocean (Section 4.1).

Additionally, users can sketch a wind deformation map, enabling simulation of natural erosion patterns caused by wind and waves (Section 4.1.1).

From these sketches, the procedural system generates a synthetic island terrain with the keep-up strategy of coral reefs (Section 4.2) and a corresponding semantic label map, where each pixel indicates its region type (island, beach, lagoon, reef, abyss) (Section 4.2.3).

User interaction

As users draw the top-view and profile-view sketches, the system provides real-time feedback on the resulting terrain. The top-view sketch influences the horizontal layout of the island, while the profile-view sketch defines its vertical structure. These sketches can be adjusted independently, allowing the user to fine-tune both the outline and elevation of the island.

While sketching the basic shape, users can apply wind deformation strokes to modify the island’s features further. These strokes

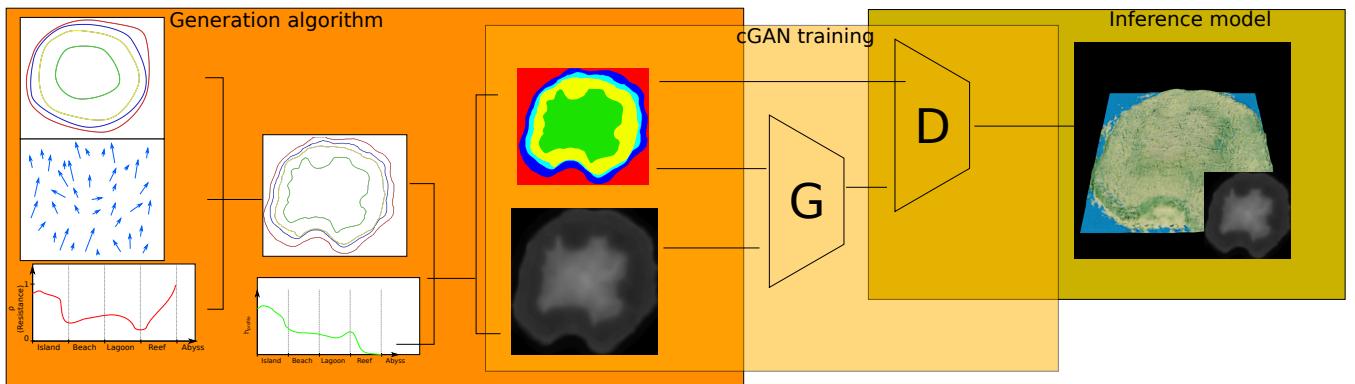


Figure 1: Our method is split in three interleaved stages: the generation process (Section 4) which creates pairs of height fields and label maps of an island from sketches, the model training (??) which use a synthetic dataset from the previous stage to obtain a cGAN model that generates height fields from label maps to remove the constraints embedded in the initial generation process, and finally, the inference process (??) uses the trained cGAN to generate the final height fields, including the coral generation process, automatically.

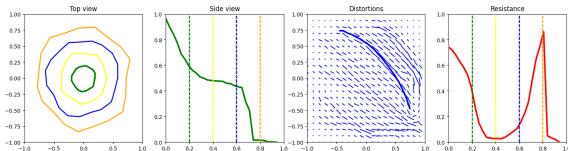


Figure 2: The user can interact directly on the island by editing the different canvases in no specific order. This UI shows, from left to right, the top-view sketch with the different outlines of each regions, the profile-view sketch with the outlines represented in dotted lines, the wind velocity sketch drawn with strokes (last stroke is visible), and the resistance function showing here a high resistance at the top of the island and on the front reef.

represent wind and wave influences, distorting the island's shape to introduce more natural, non-radial features such as indentations along the coastline, variable lagoon shapes, or concave formations. The system automatically applies these deformations, providing real-time feedback as the user interacts with the terrain.

This interactive process, combining sketches and wind deformation, allows users to quickly iterate on their designs, refining the terrain to meet specific aesthetic or functional goals.

3.2. Learning-based generation phase

We repeat this procedural generation process many times with varied parameters (different shapes, scales, subsidence levels, and wind patterns) to create a large synthetic dataset (Section 5.1). Each dataset entry consists of a label map paired with its procedurally generated terrain height field. Data augmentation is applied to the generated pairs to reduce the impact of the constraints induced from the procedural method (Section 5.2).

We use this dataset to train a Conditional Generative Adversarial Network (cGAN), specifically the pix2pix architecture, capable of

translating label semantic maps into realistic terrain height fields (??).

After training, the procedural step becomes unnecessary. To generate new island terrains, the user only needs to provide a label semantic map as input to the trained cGAN. The cGAN then synthesizes realistic island elevation details directly, capturing learned geological and geomorphological patterns from the synthetic training data (??).

User interaction

Thus, the trained cGAN provides a user-friendly interface: users draw or edit simple label maps (regions) to rapidly generate diverse, geologically plausible coral reef island terrains, incorporating realistic features such as smooth transitions between regions, detailed coral reef structures, and naturally varied shapes free from procedural constraints.

This combined procedural-and-learning approach provides a simple, flexible, and powerful tool for island terrain generation, enabling users to intuitively generate realistic and diverse coral reef islands aligned with real-world geological and biological processes such as volcanic subsidence, coral reef growth, and wind-driven erosion.

4. Procedural terrain generation

The generation of coral reef island terrains involves a structured process that takes the user's sketches and produces a complete 3D terrain model. This process begins with the creation of the initial height field based on the user's input, followed by the application of wind deformation to introduce natural variations, and concludes with the integration of coral reef features through subsidence and coral growth modeling.

The generation of coral reef islands in this system begins with two intuitive sketch-based inputs from the user: a top-view sketch and a profile-view sketch, which define the islands horizontal layout and vertical elevation profile. In addition to these sketches, the user can further refine the terrain by applying wind deformation strokes, which simulate the effects of wind and waves on the island's shape. This combination of sketches and wind inputs gives users precise control over both the island's structure and its natural variations, such as irregular coastlines or concave features. We will present the usefulness of these sketches in this section, and describe the technical details in the next section.

4.1. Initial height field generation

The top-view sketch defines the island's outline as seen from above. Using a simple drawing interface, the user delineates concentric boundaries for key regions such as the island itself, the beaches, the lagoon, and the surrounding abyss, around the canvas center. Each boundary is represented in polar coordinates, where r_p is the radial distance from the island's center and θ_p is the angular position. Allowing r to vary with θ introduces irregular, natural shapes rather than perfect circles (Figure 3).

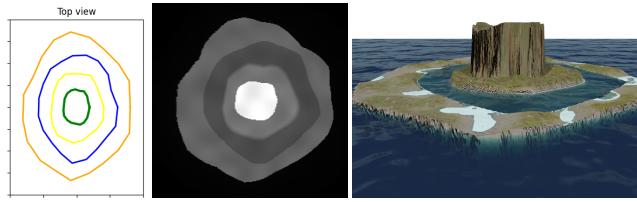


Figure 3: Using only the outlines from the top-view sketch, each point in the field is assigned a region (island, beach, lagoon, abyss), which later guides its height assignment.

The profile-view sketch defines the island's vertical elevation along any radial direction. Here, the user draws a curve that specifies height at key terrain milestones, such as the central peak, island border, beach, lagoon, and abyss, creating a continuous 1D height function $h_{\text{profile}}(\tilde{x})$ where \tilde{x} is a parametric distance measuring position along the sequence of regions. This continuous profile ensures smooth elevation transitions across all terrain features.

By combining the top-view and profile-view sketches via revolution modeling, the system generates a full 3D terrain model that matches the user's design. The process begins by transforming those two sketches into a coherent height field (Figure 5).

For any point p on the terrain, the system computes polar coordinates (r_p, θ_p) . The radial distance r_p determines which region the point belongs to (island, beach, lagoon, reef, or abyss) based on the user-defined radial limits. Those outlines from the top-view sketch provide the exact boundaries between regions.

Each point's height is computed using the profile function $h_{\text{profile}}(\tilde{x})$. Instead of using the raw radial distance r_p , we define a parametric region distance \tilde{x}_p that maps each point to a normalized position along the concentric regions (see Figure 4). The radial

space is divided by user-defined boundaries R_0, R_1, \dots, R_n , corresponding to the island center, border, beach, lagoon, and abyss.

When a point p lies between two boundaries, say R_i and R_{i+1} , its parametric distance is

$$\tilde{x}_p = i + \frac{r_p - R_i}{R_{i+1} - R_i}, \quad (1)$$

where i is the index of the region containing p (i.e., $R_i \leq r_p < R_{i+1}$). This linear mapping stretches each region's radial span to the interval $[i, i+1]$, ensuring smooth interpolation across region boundaries. The final height at point p is then

$$h(p) = h_{\text{profile}}(\tilde{x}_p). \quad (2)$$

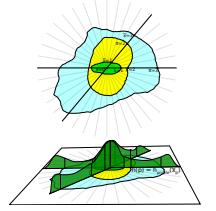


Figure 4: Parametric distance \tilde{x} depending on angle θ .

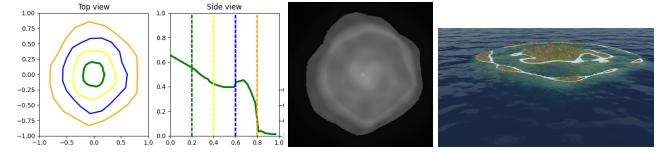


Figure 5: Given top-view and side-view outlines, the 3D result is obtained by revolution.

4.1.1. Wind deformation

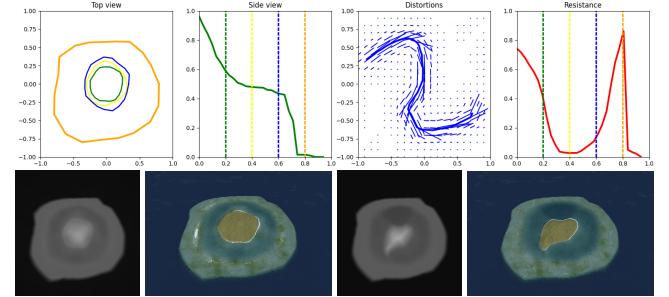


Figure 6: Defining a top-view wind vector field from user-provided strokes (top, blue) in association with a resistance function (top, red), a height field is deformed accordingly. Right: original height field and render; Left: altered results. The beach and lagoon regions are defined with low resistance, which is visible by having only these regions deformed in bottom results.

To break the radial symmetry inherent in sketch-based terrain generation and introduce more organic island shapes, we allow the user to define a wind velocity field via freehand strokes on a 2D canvas. Each stroke is represented as a parametric curve C , interpreted as a local wind flow with direction C' , strength S , and influence width σ . These strokes simulate wind and wave erosion effects on the terrain.

The deformation vector at any terrain point p is computed as a

sum over all strokes, weighted by a Gaussian falloff centered on the closest point \mathbf{p}_C^* along each curve (Figure 6, top):

$$\Phi(\mathbf{p}) = \sum_{C \in \text{curves}} S \frac{C'(\mathbf{q})}{\|C'(\mathbf{q})\|} \cdot G_\sigma (\|\mathbf{p} - \mathbf{p}_C^*\|), \quad G_\sigma(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{x^2}{2\sigma^2}}. \quad (3)$$

To preserve semantic structure across terrain regions, a resistance function $\rho(\tilde{x})$ modulates the deformation based on terrain zones such as beach, lagoon, or abyss (Figure 7). The final deformation vector becomes:

$$\tilde{\Phi}(\mathbf{p}) = (1 - \rho(\tilde{x}_{\mathbf{p}})) \cdot \Phi(\mathbf{p}). \quad (4)$$

This warp is applied to both the height field and the label map, ensuring consistent semantic deformation. For example, applying strokes to one side of a circular island creates concave coastlines while leaving high-resistance regions (e.g., the abyss) unaffected, simulating the localized impact of natural erosion (Figure 6, bottom).

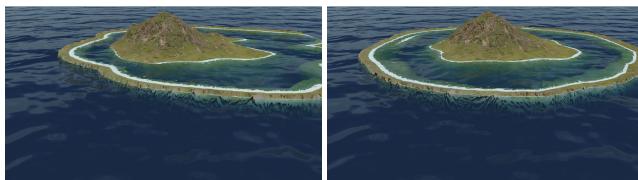


Figure 7: (Left) Given a uniform wind velocity field and a resistance function similar as ??, the coasts are smoothly eroded while the interior of the island is almost unaffected. (Right) Modifying the resistance function to affect a strong resistance to borders simulate the effect of coast reinforcements.

4.2. Coral reef modeling

Once the terrain has been generated and deformed by the wind, we simulate the long-term geological evolution of coral reef islands through two parallel processes: the subsidence of the volcanic island and the upward growth of coral reefs. As observed in nature, the volcanic base sinks over time while coral formations grow vertically to remain close to the water surface, following the “keep-up” strategy of reef development.

4.2.1. Subsidence

Subsidence is modeled by uniformly scaling the original terrain height downward, simulating the gradual sinking of the volcanic landmass due to tectonic processes. The user specifies a subsidence rate $\lambda \in [0, 1]$, which controls how much the island has sunk. The subsided terrain is computed as:

$$h_{\text{subsid}}(\mathbf{p}) = (1 - \lambda) \cdot h_0(\mathbf{p}) \quad (5)$$

This factor is applied uniformly across the island, offering a geologically plausible and computationally efficient approximation of large-scale subsidence.

4.2.2. Coral reef growth

Coral reef growth is modeled independently from the subsiding terrain. The system generates a coral-specific height field $h_{\text{coral}}(\mathbf{p})$ that remains near the sea surface regardless of the island’s vertical shift, reflecting coral growth in biologically viable depth ranges (typically 0-30 meters below sea level).

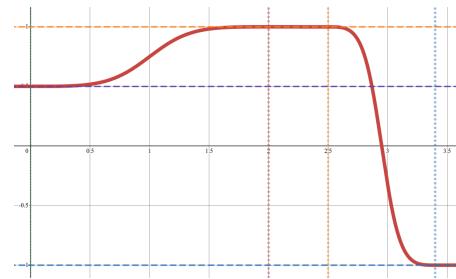


Figure 8: The modeling of the reef growth in our model is described by a piecewise function h_{coral} which is flat in the lagoon, the crest and abyss, and follows a smoothstep function as transitions for the backreef and fore reef regions. Zones’ anchor heights are represented by horizontal dashed lines; zones’ limits are dotted vertical lines.

We define distinct reef zones anchored at specific depths:

- Reef crest near sea level: $h_{\text{crest}} = -2 \text{ m}$
- Back reef and lagoon: $h_{\text{back}} = -20 \text{ m}$
- Fore reef sloping to abyss: $h_{\text{abyss}} = -100 \text{ m}$

Each reef subregion is defined over a parametric domain $x \in [0, 1]$. For instance:

- Back reef: $x_{\text{back},\text{start}} = 0, x_{\text{back},\text{end}} = 0.5$
- Reef crest: $x_{\text{crest},\text{start}} = 0.75, x_{\text{crest},\text{end}} = 0.8$
- Abyss begins at $x_{\text{abyss},\text{start}} = 1$

We model transition zones between these regions using a smoothstep operator:

$$\text{smoothstep}(x) = 3x^2 - 2x^3 \quad (6)$$

We denote the interpolating function as:

$$S(a, b, x_0, x_1, x) = a + (b - a) \text{smoothstep} \left(\frac{x - x_0}{x_1 - x_0} \right) \quad (7)$$

The complete coral height field, as displayed in Figure 8, is built as a piecewise function:

$$h_{\text{coral}}(x) = \sum_{r \in \text{subregions}} \begin{cases} h_r & \text{if } x_{r,\text{start}} \leq x \leq x_{r,\text{end}} \\ 0 & \text{otherwise} \end{cases} + \sum_{t \in \text{transitions}} \begin{cases} S(h_t, h_{t+1}, x_{t,\text{end}}, x_{t+1,\text{start}}, x) & \text{if } x_{t,\text{end}} < x < x_{t+1,\text{start}} \\ 0 & \text{otherwise} \end{cases} \quad (8)$$

4.2.3. Output

Finally, we merge the island base height field and the coral reef height field by our *ad-hoc* smooth maximum operator smax defined as:

$$\text{smax}(a, b) = \begin{cases} a + \frac{\delta}{2k} \left(\frac{1}{1+e^{-k\delta}} + \frac{1}{1-e^{-k\delta}} \right) & \text{for } a \neq b \\ a + \frac{1}{2k} & \text{for } a = b \end{cases} \quad (9)$$

with $\delta = b - a$ for conciseness, and k a sharpness parameter approximating the max operator as k grows. At $k = 5$, the operator max is already well approximated while conserving continuousness in the resulting height field $\mathcal{H}(p) = \text{smax}(h_{\text{subsidi}}(\mathbf{p}), h_{\text{coral}}(\mathbf{p}))$.

The resulting terrain represents a plausible coral reef island, where the volcanic island has subsided, and coral reefs have grown upward to keep pace with the water level. The smooth blending between the subsided terrain and the coral features ensures a natural transition between regions like the island, lagoon, and coral reefs.

One of the key strengths of this method is its flexibility as the subsidence and coral reef growth processes are modeled independently, allowing for a wide range of configurations. Users can generate plausible island terrains with or without coral features, or apply the coral reef growth simulation to existing height fields from other sources.

5. Training the cGAN

In this section, we introduce the use of a conditional Generative Adversarial Network (cGAN), specifically the pix2pix model, to enhance the island generation process by increasing the variety and flexibility of terrains. While the initial procedural algorithm can create numerous island examples, cGAN provides additional flexibility in generating more complex terrain without the rigid constraints of the procedural algorithm that stem from our initial assumptions based on coral reef formation theory.

5.1. Dataset generation

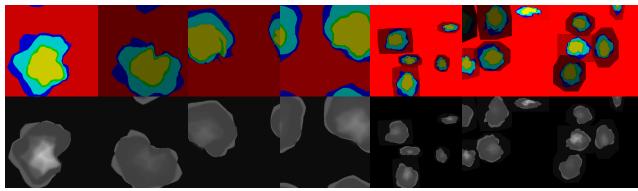


Figure 9: Using a large set of pairs of height field-label map, the training of a deep learning model result in a user-friendly interface requiring solely a hand-drawn label map to produce a 2.5D height field of the desired island.

The creation of the dataset is done through the use of the procedural algorithm for which we alter the input parameters.

For each generation, the top-view and profile-view sketches use an initial layout. Each outline of the top-view sketch is defined as a

centered circle of random radius $r_{\min} \leq r^* \leq r_{\max}$. We add another deformation based on Perlin noise such that the final contour is defined as

$$r(\theta) = r^* + \eta(\theta) \quad (10)$$

On the other hand, we define an initial profile-view sketch by defining $h_{\text{profile}}^*(\tilde{x})$ the initial height function for which fBm noise is applied to obtain

$$h_{\text{profile}}(\tilde{x}) = h_{\text{profile}}^*(\tilde{x}) \cdot \eta(\tilde{x}) \quad (11)$$

An identical process is done for the resistance function:

$$\rho(\tilde{x}) = \rho^*(\tilde{x}) \cdot \eta(\tilde{x}) \quad (12)$$

Finally, we need to generate a random wind field. The realistic nature of wind is ignored for the generation of the wind strokes in order to provide complexity and variety in the results. We generate a random number n of strokes and their path by a uniformly sampling a random number m of points. The spread and intensity of each stroke is also random.

Once all inputs are set, we generate an example for multiple level of subsidence $\lambda \in [0, 1]$ to obtain a height field incorporating the coral reef modeling and the associated label map.

The Pix2pix model was originally pretrained using RGB images. In this training phase, the images were label using the HSV (Hue, Saturation, Value) color space, where the Hue component specifically carried the label information. Both the Saturation and Value components were kept neutral, meaning they did not convey any significant label-related data. The target images, the ones the model aimed to reproduce, were formatted in RGB.

For the purpose of fine-tuning the model, we retained the use of the Hue component to encode the labels from the label map. We introduced a new dimension to the model's learning capabilities by incorporating the subsidence rate, denoted as λ , into the Value component. This addition not only utilizes the model's existing capability to interpret the HSV format but also enriches the input data, which now carries additional, valuable environmental information.

Moreover, we purposefully left the Saturation component unchanged at this stage, reserving space for potentially including another parameter in the future, which would allow us to expand the model's utility without altering the foundational HSV encoding scheme established during its initial training. By adhering to this encoding format, we ensure continuity in data representation, which maximizes the efficiency of the pretrained model. This strategic update enhances the model's adaptability and broadens its applicability to tackle new, complex challenges more effectively.

This configuration allows the process to create quickly a large quantity of data, with multiple parameters, of a single island centered in the image.

5.2. Data augmentation

To enhance the variety of the dataset and improve the model's ability to generalize, we apply several data augmentation techniques:

Translation: Since the original algorithm always centers the island, we translate the islands within the image to remove this constraint (??). This ensures that the cGAN can generate islands in any position within the frame.

Directional scaling: By scaling the terrain in one direction, we create elongated islands that resemble corridors or archipelagos, adding another layer of diversity to the dataset. Such islands are usually found on tectonic plates convergence boundaries, creating island arcs with high density of volcanic centers like the Izu-Bonin-Mariana arc system (?? shows an example of elongated island).

Copy-paste: In some cases, we combine multiple islands into a single sample, ensuring they do not overlap. The regions not covered by any island are assigned the abyss ID. Although this approach ensures non-overlapping regions, future work could explore using blending techniques to position islands more closely without the risk of overlap (??).

All augmentation techniques are applied both to the height field and the label map simultaneously to ensure consistency between the input (the label map) and the output (the height field).

6. Results and evaluation

The resulting model for coral island generation enables a high control-level from a user perspective as the unconstraint painting allows for complex scenarios while producing in real-time the resulting height fields. In this chapter we used the software Blender to provide renders directly from the outputted height fields. As our pix2pix model is trained to output 256×256 images, the resolution of the 3D models is limited by this architecture.

6.1. Control

Using deep-learning-based models, most constraints from our initial assumptions are lifted (radial layout, isolated islands, ...). The control over the overall shapes of the islands regions are given through digital painting, here using the GIMP software. Each pixel of the image are encoded in HSV, with the region identifier encoded in the Hue channel. The user may increase or decrease the subsidence level of the island by modifying the Saturation channel over the whole image (see Figure 12).

Since the model is based on statistics over the pixel values instead of hard values, users are not limited to a finite number of region identifiers, meaning that the output is more or less robust to noise (due to image compression, for example) and to the fuzzy values resulting from anti-aliasing of brushes often set by default, or resizing algorithm, by image editors, or even due to compression algorithm. The example displayed in Figure 13 presents a sketch for which the outlines of the regions are at the same time blurry and with layouts that are not expected (such as the small red regions inside the southern lagoon region or the adjacency of beach regions directly with the abyssal region) on the top figure and oversaturated on the bottom figure. The learned model does not include inconsistancies and results in plausible 3D models.

The tolerance over the input values may be used to provide even

more control about the transitions between two regions. Figure 14 shows an example of input map with regions that are leaking over neighboring regions, and the introduction of new hue values non-existent in the dataset (light green and dark green) but are the interpolated hue value of mountain regions and beach regions.

Since the procedural phase included low randomness, the output of the cGAN is limiting its unpredictability and the results to a slight change on the input create only slight changes on the output, preventing unexpected results. Figure 12 shows the result of an input map with only a variation on the subsidence level, the resulting height fields are very similar. Adding the real-time computation of outputs, it becomes possible to construct progressively a landscape and correct small mistakes to intuitively design islands inspired by real-world regions (see an reproduction of Mayotte in Figure 15).

6.2. Performances

The Python script for the initial island dataset generation is poorly optimized and takes about 2.5s per island of size 256×256 as the parallelization does not take place here. Implementing an optimized C++ version of the initial generation process reduces this execution time to 50ms per generation.

On the other hand, the inference time for a single input image of dimension 256×256 is constant whatever the complexity of the scene. Using the NVIDIA GeForce GTX 1650 Ti GPU with Python 3.10 and PyTorch version 2.5.1+cu121, the inference time measured is 5ms (std 1.1ms).

We not only show that using a neural network reduces the constraints on the generation process, but also that the execution time is only dependant on the network architecture, without influence from the dataset generation algorithm.

7. Conclusion and future work

This work has presented a novel approach to generating coral reef island terrains by combining traditional procedural methods with deep learning techniques. We first developed a procedural generation algorithm capable of creating a wide variety of island terrains through a combination of top-view and profile-view sketches, wind deformation, and subsidence and coral reef growth simulation. By applying these methods, we were able to produce realistic terrains based on geological processes, capturing key features of coral reef islands such as beaches, lagoons, and coral reefs.

To further enhance flexibility and realism in the generation process, we incorporated a Conditional Generative Adversarial Network (cGAN), using the pix2pix model to generate height maps from label maps of island features. The cGAN model allowed us to overcome some of the constraints inherent in the procedural algorithm, such as radial symmetry and fixed island positioning. With data augmentation techniques, we were able to train the cGAN on a synthetic dataset, generating varied and realistic island terrains.

Limitations While this approach brings significant advantages, there are also some limitations to consider. The reliance on a synthetic dataset means that the cGAN inherits some biases and limitations of the original procedural algorithm. This could limit the true

diversity of the terrains that the model can generate, as the output is confined by the patterns present in the training data. Additionally, the cGAN model's internal logic lacks transparency, offering limited user control over the generation process once the model has been trained. This contrasts with traditional procedural methods, which typically allow for real-time tweaking of parameters.

Future work Further improvements could be made to the synthetic dataset. Incorporating more complex geological processes, such as wave erosion or tidal influences, could lead to even more realistic terrains. Additionally, refining the way islands are blended in multi-island samples, or adding more diverse input conditions (e.g., different geological settings), could help the model generalize better and produce more varied and dynamic landscapes. While the current model allows for rapid terrain generation, adding more options for users to interact with the cGAN, such as tweaking parameters like wind strength or island size, could enhance the flexibility of the system. Many other neural networks models could be exploited to increase the possibilities, such as newer variants of cGANs Park et al., 2019, or models with style transfer functionalities Gatys et al., 2015; Zhu et al., 2020 in order to change the overall aspect of a terrain Perche et al., 2023b, 2023a, use text-to-images models Rombach et al., 2021; Radford et al., 2021 to generate height fields from a verbal prompt, or super-resolution models Dong et al., 2014 to increase the definition of details in the final output Guérin, Digne, Galin, and Peytavie, 2016.

References

- Beckham, C., & Pal, C. (2017). A step towards procedural terrain generation with gans. <http://arxiv.org/abs/1707.03383> (cit. on p. 2).
- Beneš, B., Těšínský, V., Hornýš, J., & Bhatia, S. K. (2006). Hydraulic erosion. *Computer Animation and Virtual Worlds*, 17, 99–108. <https://doi.org/10.1002/cav.77> (cit. on p. 2).
- Cordonnier, G., Braun, J., Cani, M.-P., Benes, B., Galin, É., Peytavie, A., & Guérin, É. (2016). Large scale terrain generation from tectonic uplift and fluvial erosion. *Computer Graphics Forum*, 35, 165–175. <https://doi.org/10.1111/cgf.12820> (cit. on p. 2).
- Cordonnier, G., Cani, M.-P., Beneš, B., Braun, J., & Galin, É. (2017). Sculpting mountains: Interactive terrain modeling based on subsurface geology. *IEEE Transactions on Visualization and Computer Graphics*, 24. <https://doi.org/10.1109/TVCG.2017.2689022> (cit. on p. 2).
- Cordonnier, G., Galin, É., Gain, J., Beneš, B., Guérin, É., Peytavie, A., & Cani, M.-P. (2017). Authoring landscapes by combining ecosystem and terrain erosion simulation. *ACM Transactions on Graphics*, 36. <https://doi.org/10.1145/3072959.3073667> (cit. on p. 2).
- Dong, C., Loy, C. C., He, K., & Tang, X. (2014). Image super-resolution using deep convolutional networks. <http://arxiv.org/abs/1501.00092> (cit. on p. 8).
- Ebert, D. S., Peachey, D., Worley, S., Hart, J. C., Musgrave, K., Perlin, K., & Mark, W. R. (2003). *Texturing and modeling, a procedural approach* (Vol. Third edition). Morgan Kaufmann. (Cit. on p. 2).
- Ecormier-Nocca, P., Cordonnier, G., Carrez, P., Moigne, A. M., Memari, P., Benes, B., & Cani, M. P. (2021). Authoring consistent landscapes with flora and fauna. *ACM Transactions on Graphics*, 40. <https://doi.org/10.1145/3450626.3459952> (cit. on p. 2).
- Fournier, A., Fussell, D., & Carpenter, L. (1982). Computer rendering of stochastic models. *Communications of the ACM*, 25, 371–384. <https://doi.org/10.1145/358523.358553> (cit. on p. 2).
- Gain, J., Marais, P., & Straßer, W. (2009). Terrain sketching. *Proceedings of I3D 2009: The 2009 ACM SIGGRAPH Symposium on Interactive 3D Graphics and Games*, 1, 31–38. <https://doi.org/10.1145/1507149.1507155> (cit. on p. 2).
- Galin, E., Guérin, É., Peytavie, A., Cordonnier, G., Cani, M.-P., Benes, B., & Gain, J. (2019). A review of digital terrain modeling. *Computer Graphics Forum*, 38, 553–577. <https://doi.org/10.1111/cgf.13657> (cit. on p. 2).
- Gasch, C., Chover, M., Remolar, I., & Rebollo, C. (2020). Procedural modelling of terrains with constraints. *Multimedia Tools and Applications*, 79, 31125–31146. <https://doi.org/10.1007/s11042-020-09476-3> (cit. on p. 2).
- Gatys, L. A., Ecker, A. S., & Bethge, M. (2015). A neural algorithm of artistic style. <http://arxiv.org/abs/1508.06576> (cit. on p. 8).
- Génevaux, J.-D., Galin, É., Peytavie, A., Guérin, É., Briquet, C., Grosbellet, F., & Beneš, B. (2015). Terrain modelling from feature primitives. <https://doi.org/https://doi.org/10.1111/cgf.12530> (cit. on p. 2).
- Guérin, E., Peytavie, A., Masnou, S., Digne, J., Sauvage, B., Gain, J., & Galin, E. (2022). Gradient terrain authoring. *Computer Graphics Forum*, 41, 85–95. <https://doi.org/10.1111/cgf.14460> (cit. on p. 2).
- Guérin, É., Digne, J., Galin, É., & Peytavie, A. (2016). Sparse representation of terrains for procedural modeling. *Computer Graphics Forum*, 35, 177–187. <https://doi.org/10.1111/cgf.12821> (cit. on p. 8).
- Guérin, É., Digne, J., Galin, É., Peytavie, A., Wolf, C., Beneš, B., & Martinez, B. (2017). Interactive example-based terrain authoring with conditional generative adversarial networks. *ACM Transactions on Graphics*, 36. <https://doi.org/10.1145/3130800.3130804> (cit. on p. 2).
- Hnaidi, H., Guérin, E., Akkouche, S., Peytavie, A., & Galin, E. (2010). Feature based terrain generation using diffusion equation. *Computer Graphics Forum*, 29, 2179–2186. <https://doi.org/10.1111/j.1467-8659.2010.01806.x> (cit. on p. 2).
- Hopley, D. (2014). *Encyclopedia of modern coral reefs : Structure, form and process*. Springer, Credo Reference. (Cit. on p. 1).
- Isola, P., Zhu, J.-Y., Zhou, T., & Efros, A. A. (2017). Image-to-image translation with conditional adversarial networks. *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 5967–5976. <https://doi.org/10.1109/CVPR.2017.632> (cit. on p. 2).
- Körner, C., Spehn, E. M., Bugmann, H., & Zurich, E. (2014). *Mountain systems* (tech. rep.). <https://www.researchgate.net/publication/238663492> (cit. on p. 1).

- Lyons, M. B., Murray, N. J., Kennedy, E. V., Kovacs, E. M., Castro-Sanguino, C., Phinn, S. R., Acevedo, R. B., Alvarez, A. O., Say, C., Tudman, P., Markey, K., Roe, M., Canto, R. F., Fox, H. E., Bambic, B., Lieb, Z., Asner, G. P., Martin, P. M., Knapp, D. E., ... Roelfsema, C. M. (2024). New global area estimates for coral reefs from high-resolution mapping. *Cell Reports Sustainability*, 1, 100015. <https://doi.org/10.1016/j.crsus.2024.100015> (cit. on p. 1).
- Mei, X., Decaudin, P., & Hu, B. G. (2007). Fast hydraulic erosion simulation and visualization on gpu. *Proceedings - Pacific Conference on Computer Graphics and Applications*, 47–56. <https://doi.org/10.1109/PG.2007.27> (cit. on p. 2).
- Musgrave, F. K., Kolb, C. E., & Mace, R. S. (1989). The synthesis and rendering of eroded fractal terrains. *Proceedings of the 16th Annual Conference on Computer Graphics and Interactive Techniques, SIGGRAPH 1989*, 41–50. <https://doi.org/10.1145/74333.74337> (cit. on p. 2).
- Natali, M., Viola, I., & Patel, D. (2012). Rapid visualization of geological concepts. *Brazilian Symposium of Computer Graphic and Image Processing*, 150–157. <https://doi.org/10.1109/SIBGRAPI.2012.29> (cit. on p. 2).
- Neidhold, B., Wacker, M., & Deussen, O. (2005). Interactive physically based fluid and erosion simulation. *Natural Phenomena*, 25–32. <http://graphics.uni-konstanz.de/publikationen/Neidhold2005InteractivePhysicallyBased/Neidhold2005InteractivePhysicallyBased.pdf> (cit. on p. 2).
- Olsen, J. (2004). Realtime procedural terrain generation. *Department of Mathematics And Computer Science*, 20. <https://pdfs.semanticscholar.org/5961/c577478f21707dad53905362e0ec4e6ec644.pdf> %5Cnhttp://www.tsi.enst.fr/~bloch/P6/PRREC/terrain_generation.pdf %5Cn<http://web.mit.edu/cesium/Public/terrain.pdf> (cit. on p. 2).
- Panagiotou, E., & Charou, E. (2020). Procedural 3d terrain generation using generative adversarial networks. [http://arxiv.org/abs/2010.06411](https://arxiv.org/abs/2010.06411) (cit. on p. 2).
- Park, J., Redwine, J., Hill, T. D., & Kotun, K. (2019). Water resource and ecotone transformation in coastal ecosystems. *Ecological Modelling*, 405, 69–85. <https://doi.org/10.1016/j.ecolmodel.2019.04.015> (cit. on p. 8).
- Perche, S., Peytavie, A., Benes, B., Galin, E., & Guérin, E. (2023a). Styledem: A versatile model for authoring terrains. [http://arxiv.org/abs/2304.09626](https://arxiv.org/abs/2304.09626) (cit. on p. 8).
- Perche, S., Peytavie, A., Benes, B., Galin, E., & Guérin, E. (2023b). Authoring terrains with spatialised style. *Computer Graphics Forum*, 42. <https://doi.org/10.1111/cgf.14936> (cit. on p. 8).
- Perlin, K. (1985). An image synthesizer. *ACM SIGGRAPH Computer Graphics*, 19, 287–296. <https://doi.org/10.1145/325165.325247> (cit. on p. 2).
- Perlin, K. (2001). Noise hardware. *Real-Time Shading SIGGRAPH Course Notes* (cit. on p. 2).
- Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, A., Mishkin, P., Clark, J., Krueger, G., & Sutskever, I. (2021). Learning transferable visual models from natural language supervision. <http://arxiv.org/abs/2103.00020> (cit. on p. 8).
- Rombach, R., Blattmann, A., Lorenz, D., Esser, P., & Ommer, B. (2021). High-resolution image synthesis with latent diffusion models. <http://arxiv.org/abs/2112.10752> (cit. on p. 8).
- Schott, H., Paris, A., Fournier, L., Guérin, E., & Galin, E. (2023). Large-scale terrain authoring through interactive erosion simulation. *ACM Transactions on Graphics*, 42, 1–15. <https://doi.org/10.1145/3592787> (cit. on p. 2).
- Sisodia, Y. (2022). Gan-generated terrain for game assets. *Indian Journal of Artificial Intelligence and Neural Networking*, 2, 1–3. <https://doi.org/10.54105/ijainn.F1060.102622> (cit. on p. 2).
- Smelik, R. M., Kraker, K. J. D., Groenewegen, S. A., Tutenel, T., & Bidarra, R. (2009). A survey of procedural methods for terrain modelling. *Proceedings of the CASA workshop on 3D advanced media in gaming and simulation (3AMIGAS)* (cit. on p. 2).
- Spick, R., & Walker, J. A. (2019). Realistic and textured terrain generation using gans. *Proceedings - CVMP 2019: 16th ACM SIGGRAPH European Conference on Visual Media Production*. <https://doi.org/10.1145/3359998.3369407> (cit. on p. 2).
- Talgorn, F. X., & Belhadj, F. (2018). Real-time sketch-based terrain generation. *ACM International Conference Proceeding Series*, 13–18. <https://doi.org/10.1145/3208159.3208184> (cit. on p. 2).
- Voulgaris, G., Mademlis, I., & Pitas, I. (2021). Procedural terrain generation using generative adversarial networks. *European Signal Processing Conference, 2021-August*, 686–690. <https://doi.org/10.23919/EUSIPCO54536.2021.9616151> (cit. on p. 2).
- Wulff-Jensen, A., Rant, N. N., Møller, T. N., & Billeskov, J. A. (2018, January). Deep convolutional generative adversarial network for procedural 3d landscape generation based on dem. In *Interactivity, game creation, design, learning, and innovation* (pp. 85–94). Springer. https://doi.org/10.1007/978-3-319-76908-0_9 (cit. on p. 2).
- Zhu, D., Cheng, X., Zhang, F., Yao, X., Gao, Y., & Liu, Y. (2020). Spatial interpolation using conditional generative adversarial neural networks. *International Journal of Geographical Information Science*, 34, 735–758. <https://doi.org/10.1080/13658816.2019.1599122> (cit. on p. 8).

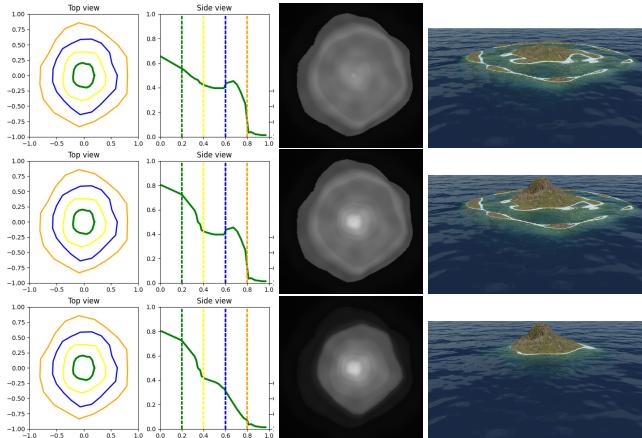


Figure 10: Providing a smooth function between each region results in islands with plausible reliefs. We fixed the outlines while editing only the height function in order to produce, from top to bottom, a low island, a coral reef island, and finally an identical island without the reef.

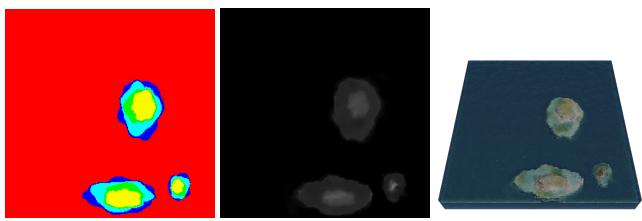


Figure 11: By applying our three data augmentation functions, the deep learning model learns to overcome some constraints previously set by the initial algorithm: (A) the translation removes the constraint to have an island ultimately at the center of the map, (B) the directional scaling, typical from image processing, reduces the symmetry constraint on the results and (C) the copy-paste unlock the possibility to obtain more than one island per map.

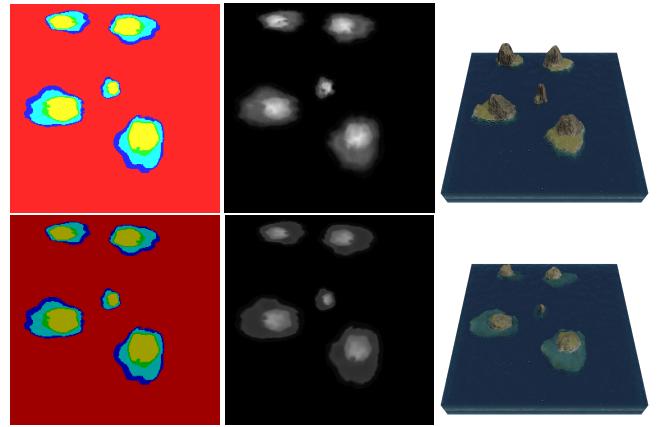


Figure 12: An identical label map yield similar height fields over multiple inferences from the model, even after modifying the subsidence factor (visible in the luminosity of the input image).

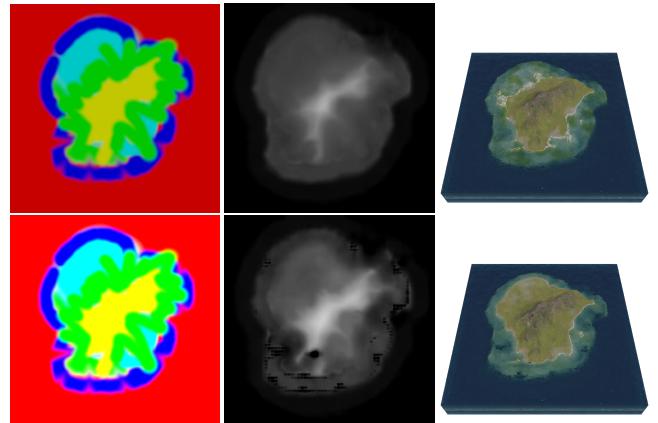


Figure 13: Using a generative neural network allows a higher level of tolerance on the user input. Here the user used a fuzzy brush to draw the label map, resulting in some pixels that are inconsistent with the dataset and unlogical island layouts (some small "abyss" regions [red] are found between "beach" [green], "lagoon" [cyan] and "reef" [blue]). The model ignores the inconsistencies even for over-saturated pixels.

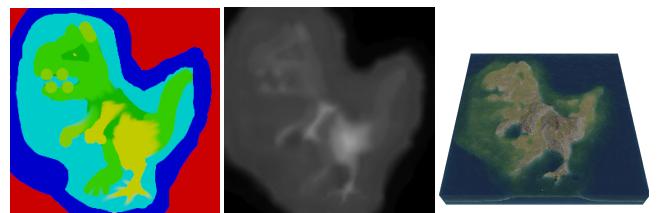


Figure 14: Without constraints on the generation, the user may use unrealistic layout and the neural network will however output a plausible result.

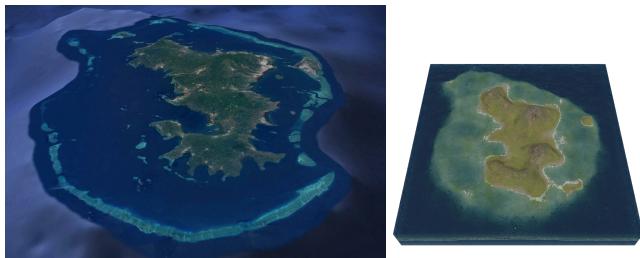


Figure 15: Comparison between real (left) and synthetic (right) islands of Mayotte.

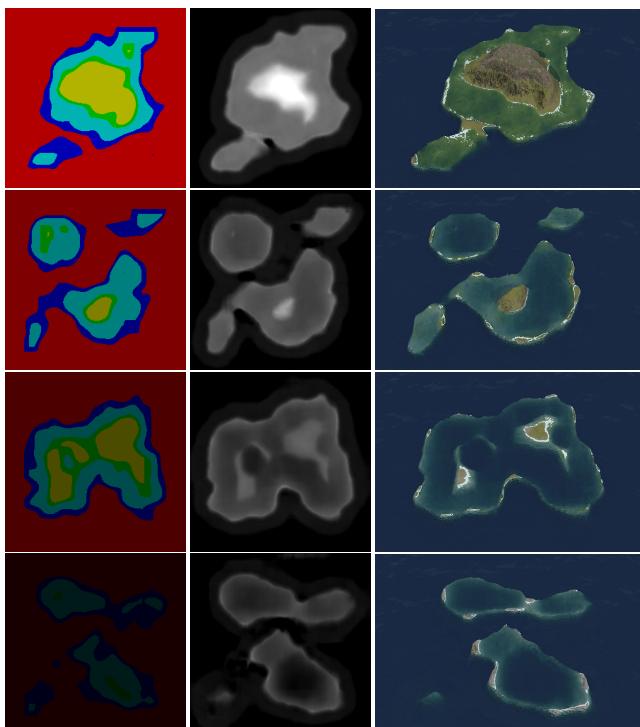


Figure 16: Starting from random Perlin noise, transformed into a label map, we can generate a large variety of results.