

Procedural and learning-based generation of coral reef islands

Submission1265



Figure 1: Given a free hand-drawn sketch of the different regions of an island, our method generate a corresponding height field using neural networks. Subsidence (gradual sinking of land) is controlled by the user in the luminosity channel of the input.

Abstract

We propose a procedural method for generating single volcanic islands with coral reefs using user sketching from two projections: a top view, which defines the island’s shape, and a profile view, which outlines its elevation. These projections, commonly used in geological and remote sensing domains, are complemented by a user-defined wind field, applied as a distortion field to deform the island’s shape, mimicking the effects of wind and waves on the long term and enabling finer user control. We then model the growth of coral on the island and its surroundings to construct the reef following biological observations. Based on these inputs, our method generates a height field of the island. Our method creates a large variety of island models to compose a dataset used for training a conditional Generative Adversarial Network (cGAN). By applying data augmentation, the cGAN allows for even greater variety in the generated islands, providing users with higher freedom and intuitive controls over the shape and structure of the final output.

1 **Keywords:** Procedural modeling, Terrain generation, cGAN,
2 Coral reef, Sketch-based interface

3 1. Introduction

4 Simulating the formation of coral reef islands presents significant
5 challenges due to the complex interplay of geological, environmental-
6 tal, and biological factors (Hopley, 2014). One major difficulty lies
7 in capturing the long-term subsidence of volcanic islands, which
8 occurs over millions of years, while concurrently simulating the
9 upward growth of coral reefs that rely on environmental conditions
10 such as water depth, temperature, and sunlight. This combination
11 of slow geological processes and dynamic biological growth is dif-
12 ficult to replicate accurately in a computational model.

13 The biological aspects of coral growth are inherently tied to envi-
14 ronmental factors. Coral reefs grow only within a specific range of
15 water depth and sunlight, and their growth patterns are affected by
16 the health of the reef ecosystem and the availability of resources.
17 Accurately modeling these biological dependencies in a procedu-
18 ral system is complex, since these factors are numerous and diffi-

19 cult to generalize. Additionally, the scarcity of data available ob-
20 structes the global understanding of these ecosystems. In a recent
21 high-resolution mapping of shallow coral reefs (Lyons et al., 2024),
22 researchers estimated the total surface area of this biome to cover
23 less than 0.7% of Earth’s area, and more specifically, coral habitats
24 constitute less than 0.2%.

25 Moreover, the need to design plausible reef structures in digi-
26 tal environments for applications such as scientific visualization,
27 simulation, or digital entertainment, calls for generative algorithms
28 that balance realism with controllability. Existing terrain generation
29 methods, such as Perlin noise-based algorithms or uplift-erosion
30 models, are often ill-suited for these processes. While they can gen-
31 erate natural-looking landscapes (such as alpine landscapes, repre-
32 senting about a quarter of land area (Körner et al., 2014)), they do
33 not account for the unique geological and biological interactions
34 that govern coral reef island formation, thus missing coherence.
35 Capturing these dynamics, while also providing user control during
36 the modeling of a terrain, requires a balance between plausibility
37 and flexibility, allowing for both accurate, though computationally

38 expensive, simulation of natural processes and intuitive user control
 39 in interactive time.

40 Despite advances in terrain generation, existing methods struggle
 41 with user-controlled design of specific island shapes and achieving
 42 realism without real data. Coral reef islands exemplify this gap:
 43 we lack datasets to directly train deep learning models, and purely
 44 procedural methods require expert tuning to mimic their features.

45 To address these issues, we propose to use curve-based modeling
 46 techniques as an initial step in our approach as a means to
 47 efficiently create a large and diverse set of training examples for
 48 a learning-based model. Each synthetic example is represented by
 49 a terrain height field and a corresponding semantic label map that
 50 marks different regions, providing structured input-output pairs for
 51 the learning stage.

52 We trained a conditional Generative Adversarial Network
 53 (cGAN) as the core of our learning-based approach (Mirza and
 54 Osindero, 2014; Isola et al., 2017). A cGAN is a type of deep learning
 55 model that learns to generate realistic data based on an input
 56 condition or context. In our case, the cGAN takes as input the
 57 semantic label map of an island generated by the procedural step and
 58 learns to produce a plausible island height field that matches this
 59 layout. By training on a large variety of examples from our curve-
 60 based algorithm, the cGAN captures the subtle terrain features and
 61 variations characteristic of coral reef islands. This approach sur-
 62 passes the capabilities of traditional procedural rules, aided by ex-
 63 tensive data augmentation. The cGAN can be used on its own to
 64 generate new island terrains with simplified and more intuitive user
 65 inputs through digital drawing, and the model will generate a real-
 66 stic island terrain accordingly.

67 The key contributions are as follows: 1) a novel curve-based pro-
 68 cedural algorithm for shaping island terrains from top and profile
 69 views, 2) the training of a deep learning model on synthetic data
 70 derived from procedural rules, serving as an abstraction layer that
 71 hides underlying complexity, 3) a demonstration that the cGAN ap-
 72 proach tolerates imprecise, low-detail user input sketches, broadening
 73 usability, without the need for cutting-edge network architec-
 74 tures, and 4) an insight that procedural generation remains essential
 75 to produce training data in data-sparse domains such as coral reef
 76 islands. These contributions collectively show a pathway to blend
 77 user-driven design with learning-based generation in terrain mod-
 78 eling, especially for data-sparse domains.

79 2. Related work

80 Procedural terrain generation spans a spectrum from purely noise-
 81 driven algorithms to physics-based simulations and, more recently,
 82 data-driven methods. In the context of coral reef islands, where
 83 both long-term geological subsidence and biogenic reef accretion
 84 interplay, existing approaches fall short in one of two ways: they
 85 either lack ecological or geological grounding, or they offer insuf-
 86 ficient authoring control.

87 **Procedural and simulation-based methods** Noise-based tech-
 88 niques such as Perlin noise (Perlin, 1985), Simplex noise (Perlin,
 89 2001), and the Diamond-Square algorithm (Fournier et al., 1982)
 90 (often extended via fBm or multifractal noise (Musgrave et al.,

91 1989; Ebert et al., 2003)) remain popular for their simplicity and
 92 speed. Island shapes are generated by modulating noise with radial
 93 falloff masks (Olsen, 2004), but these methods cannot reproduce
 94 reef rings, lagoons, or atoll structures in a geologically coherent
 95 manner. They treat terrain purely as a signal-processing problem,
 96 separated from processes like volcanic subsidence or coral growth
 97 (Smelik et al., 2009; Galin et al., 2019).

98 Simulation-based approaches introduce surface deformations by
 99 modeling erosion (Beneš et al., 2006; Neidhold et al., 2005; Mei et
 100 al., 2007), tectonic uplift (Cordonnier, Braun, et al., 2016; Cordonnier,
 101 Cani, et al., 2017; Schott et al., 2023), or vegetation-terrain
 102 feedback (Ecormier-Nocca et al., 2021; Cordonnier, Galin, et al.,
 103 2017). Hydraulic and thermal erosion capture fluvial networks and
 104 slope-driven mass wasting, but they omit underwater sedimentation
 105 and biogenic carbonate accretion. Tectonic and isostatic models ex-
 106 cel at orogeny but ignore coral reef dynamics, while vegetation-
 107 based methods do not generalize to marine ecosystems. Conse-
 108 quently, none of these simulations jointly capture the slow sub-
 109 sidence of volcanic islands and the compensatory growth of sur-
 110 rounding reefs on the timescales required for atoll formation.

111 **Sketch-based terrain modeling** Sketch-driven interfaces bridge
 112 user intent and procedural detail. Curve-based systems let users
 113 draw ridges, valleys, or coastlines that guide surface deformation
 114 and noise propagation (Gain et al., 2009; Hnaidi et al., 2010).
 115 Constraint-based methods extend this by enforcing absolute eleva-
 116 tion or slope values at control curves or points, solved via diffusion
 117 or fractal interpolation (Gasch et al., 2020; Talgorn and Belhadj,
 118 2018), and even gradient-domain editing for slope control (Guérin,
 119 Peytavie, et al., 2022). Semantic approaches encode high-level "ter-
 120 rain atoms" from a dictionary of primitives (Génevaux et al., 2015)
 121 or interpret geological schematics into 3D models (Natali et al.,
 122 2012). While these techniques grant artists fine-grained control,
 123 they typically lack ecological constraints and have not been tailored
 124 to marine biogeomorphology.

125 **Learning-based terrain synthesis** Deep generative models offer
 126 a way to learn complex patterns without explicit procedural rules.
 127 Unconditional GANs have been applied to digital elevation maps of
 128 mountains (Wulff-Jensen et al., 2018) and joint height-texture syn-
 129 thesis (Spick and Walker, 2019), but their reliance on latent noise
 130 prevents precise layout control. Two-stage pipelines use an initial
 131 GAN for heightmaps and a conditional GAN for textures (Beckham
 132 and Pal, 2017) or reverse the order (imagery-to-DEM) (Panagiotou
 133 and Charou, 2020), yet still lack the ability to guide authoring.

134 Conditional GANs (cGANs) extend image-to-image translation
 135 methods such as pix2pix (Isola et al., 2017) to terrain, enabling
 136 sketch- or label map-conditioned generation. Prior work includes
 137 sketch-to-DEM translation for generic landforms (Guérin, Digne,
 138 Galin, Peytavie, et al., 2017) and sparse "altitude dot" conditioning
 139 (Voulgaris et al., 2021), as well as cGANs that invert satellite im-
 140 agery into elevation (Sisodia, 2022). However, these models require
 141 extensive paired real-world datasets which are scarce for coral reef
 142 islands.

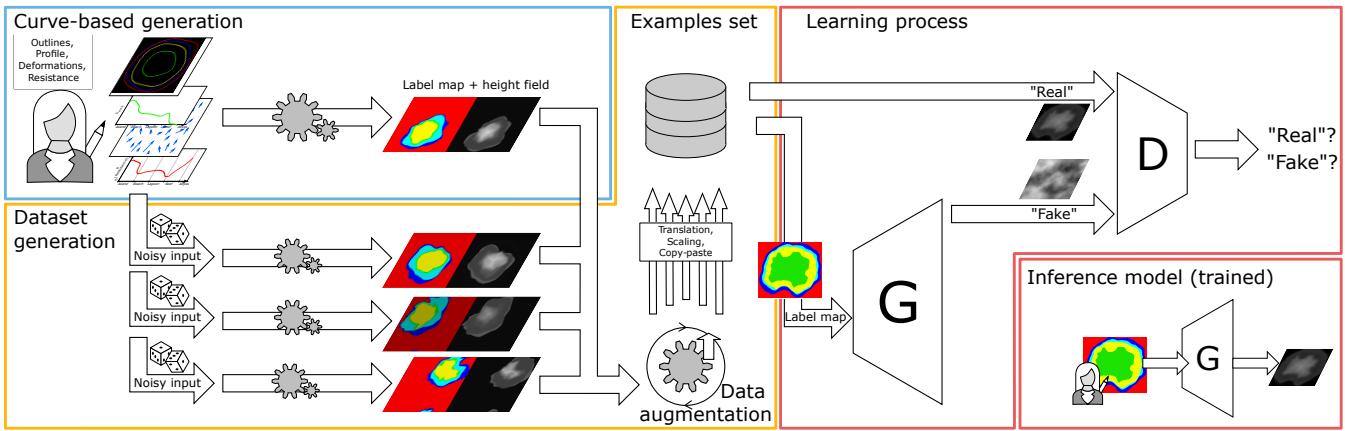


Figure 2: Our pipeline first prompts the user to create single pairs of coral reef island height fields and label maps using our proposed curve-based modeling algorithm (Section 4). The same algorithm is applied multiple times on randomly altered versions of the user input to create a dataset, which we further enhance with data augmentation techniques (Section 5). Finally, we train a conditional Generative Adversarial Network, which can then be used standalone to create height fields from labelled sketches (Section 6).

143 3. Overview

144 Our method for procedural generation of coral reef islands is composed
 145 of two independent modeling techniques shown in Figure 2.
 146 A first curve-based algorithm (top-left blue block, presented in
 147 Section 4), parametrizes the surface of islands by a top-view sketch
 148 interface, describing the outlines of its constituent regions and a
 149 stroke-based wind field deforming them, as well as a profile-view
 150 sketch describing for each region its altitude and resistance to
 151 deformations. User inputs are given as 1D functions (altitude and re-
 152 sistance), a 1D polar function (island outlines) and a 2D vector field
 153 (wind field) as shown in Figure 3, which are convinient to randomize
 154 through the use of noise, but limit the user's freedom to model
 155 arbitrary shapes.

156 We propose a second learning-based generation algorithm (right
 157 red blocks Figure 2, developed in Section 6), which trains a neu-
 158 ral generator G to transform a labelled image into a height field.
 159 Parallelly, a discriminator D is trained to distinguish height fields
 160 provided from the training set against height fields generated by
 161 G . This adversarial training strengthens the outputs of the genera-
 162 tor, up to the point where we discard the discriminator and train-
 163 ing data, only keeping the generation step (bottom-right). Users are
 164 then able to provide unseen label maps and obtain height fields as
 165 close as possible to the training dataset. This method for terrain
 166 modeling is poorly constrained, however requires a large amount
 167 of prior data in its training set, which is not available in our case.

168 We connect these two algorithms through a process of dataset
 169 generation (central orange block Figure 2, described in Section 5)
 170 in order to enforce the benefits of each while reducing their limita-
 171 tions. Given the initial curve-based user inputs, we create a dataset
 172 composed of similar samples processed by our first algorithm, raster-
 173 ized into a 2D image of the parametrized regions, and paired
 174 with the resulting height field. We then greatly increase the size
 175 of the dataset by employing various data augmentation techniques
 176 such as translations, scaling and copy-pasting of multiple islands in

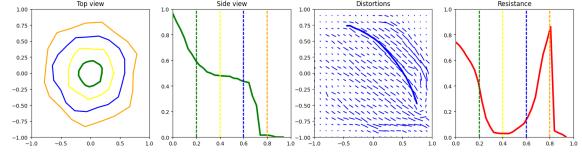


Figure 3: The user can interact directly on the island by editing the different canvases in any order. This UI shows, from left to right, the top-view sketch with the different outlines of each regions, the profile-view sketch with the outlines represented by dotted lines, the wind velocity sketch drawn with strokes (last stroke is visible), and the resistance function showing here a high resistance at the top of the island and on the front reef.

177 a single sample. We then obtain a dataset large enough for training
 178 our cGAN and enable users to procedurally create coral reef islands
 179 with a high level of control.

180 4. Curve-based island generation

181 The generation of coral reef island terrains involves a structured
 182 process that takes the user's sketches and produces a complete 3D
 183 terrain model. This process begins with the creation of the initial
 184 height field based on the user's input, followed by the application
 185 of wind deformation to introduce natural variations, and concludes
 186 with the integration of coral reef features through subsidence and
 187 coral growth modeling.

188 The generation of coral reef islands in our system begins with
 189 two intuitive curve-based inputs from the user: a top-view sketch
 190 and a profile-view sketch, which define the island's horizontal lay-
 191 out and vertical elevation profile. In addition to these sketches, the
 192 user can further refine the terrain by applying wind deformation
 193 strokes, which simulate the effects of wind and waves on the is-
 194 land's shape. This combination of sketches and wind inputs gives

users precise control over both the islands structure, and its natural variations, such as irregular coastlines or concave features. We present the usefulness of these sketches in this section, and describe the technical details in the next section.

4.1. Initial height field generation

The top-view sketch defines the island's outline as seen from above. Using a simple drawing interface, the user delineates concentric boundaries for key regions such as the island itself, the beaches, the lagoon, and the surrounding abyss, around the center of the canvas. Each boundary is represented in polar coordinates, where r_p is the radial distance from the island's center and θ_p is the angular position. Allowing r to vary with θ introduces irregular, natural shapes rather than perfect circles (Figure 4).

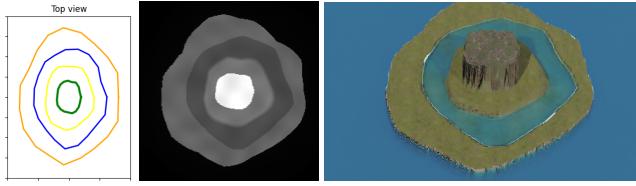


Figure 4: Using only the outlines from the top-view sketch, each point in the field is assigned a region (island, beach, lagoon, reef, abyss), which later guides its height assignment.

The profile-view sketch defines the island's vertical elevation along any radial direction. Here, the user draws a curve that specifies height at key terrain milestones, such as the central peak, island border, beach, lagoon, reef, and abyss, creating a continuous 1D height function $h_{\text{profile}}(\tilde{x})$, where \tilde{x} is a parametric distance measuring position along the sequence of regions. This continuous profile ensures smooth elevation transitions across all terrain features.

By combining the top-view and profile-view sketches via revolution modeling, the system generates a full 3D terrain model that matches the user's design. The process begins by transforming these two sketches into a coherent height field (Figure 6).

For any point p on the terrain, the system computes polar coordinates (r_p, θ_p) . The radial distance r_p determines which region the point belongs to (island, beach, lagoon, reef, or abyss) based on the user-defined radial limits. Those outlines from the top-view sketch provide the exact boundaries between regions.

Each point's height is computed using the profile function $h_{\text{profile}}(\tilde{x})$. Instead of using the raw radial distance r_p , we define a parametric region distance \tilde{x}_p that maps each point to a normalized position along the concentric regions (see Figure 5). The radial space is divided by user-defined boundaries R_0, R_1, \dots, R_n , corresponding to the island center, border, beach, lagoon, reef, and abyss.

When a point p lies between two boundaries, say R_i and R_{i+1} , its parametric distance is

$$\tilde{x}_p = i + \frac{r_p - R_i}{R_{i+1} - R_i}, \quad (1)$$

where i is the index of the region containing p (i.e., $R_i \leq r_p < R_{i+1}$). This linear mapping stretches each region's radial span to the interval $[i, i+1]$, ensuring smooth interpolation across region boundaries. The final height at point p is then $h(p) = h_{\text{profile}}(\tilde{x}_p)$.

4.1.1. Wind deformation

To break the radial symmetry inherent to our curve-based terrain generation and introduce more organic island shapes, we allow the user to define a wind velocity field via freehand strokes on a 2D canvas. Each stroke is represented as a parametric curve C , interpreted as a local wind flow with direction given by the derivative C' , strength S , and influence width σ . These strokes simulate wind and wave erosion effects on the terrain.

The deformation vector at any terrain point p is computed as a sum over all strokes:

$$\Phi(p) = \sum_{C \in \text{curves}} S \frac{C'(\mathbf{q})}{\|C'(\mathbf{q})\|} \cdot G_\sigma(\|\mathbf{p} - \mathbf{p}_C^*\|) \quad (2)$$

weighted by a Gaussian falloff centered on the closest point \mathbf{p}_C^* along each curve (Figure 7, top):

$$G_\sigma(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{x^2}{2\sigma^2}}. \quad (3)$$

To preserve semantic structure across terrain regions, a resistance function $\rho(\tilde{x})$ modulates the deformation based on terrain zones such as beach, lagoon, reef, or abyss (Figure 8). The final deformation vector becomes:

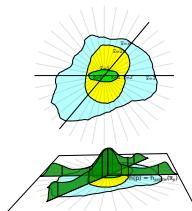


Figure 5: Parametric distance \tilde{x} depending on angle θ .

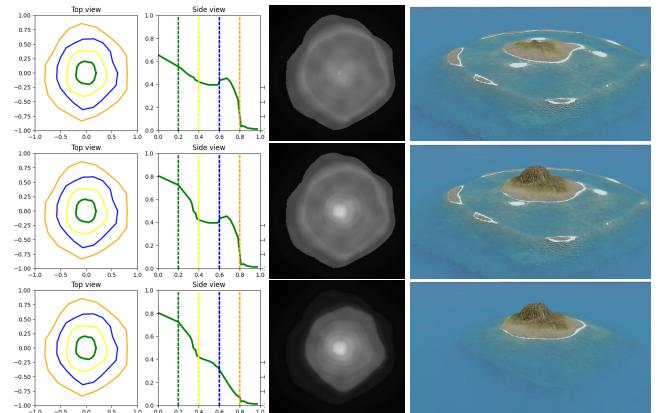


Figure 6: Providing a smooth function between each region results in islands with plausible reliefs. We fixed the outlines while editing only the height function in order to produce, from top to bottom, a low island, a coral reef island, and finally an identical island without the reef.

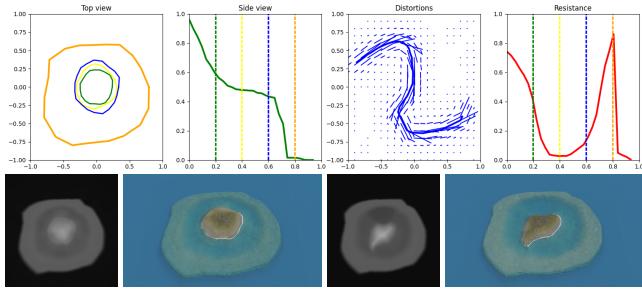


Figure 7: A top-view wind vector field is defined from user-provided strokes (top, blue), in association with a resistance function (top, red); a height field is deformed accordingly. Left: original height field and render; right: altered results. The beach and lagoon regions are defined with low resistance, which is visible by having only these regions deformed in bottom results.

landmass due to tectonic processes. The user specifies a subsidence rate $\lambda \in [0, 1]$, which controls how much the island has sunk. The subsided terrain is computed as:

$$h_{\text{subsid}}(\mathbf{p}) = (1 - \lambda) \cdot h_0(\mathbf{p}). \quad (5)$$

This factor is applied uniformly across the island, offering a geologically plausible and computationally efficient approximation of large-scale subsidence.

4.2.2. Coral reef growth

Coral reef growth is modeled independently from the subsiding terrain. The system generates a coral-specific height field $h_{\text{coral}}(\mathbf{p})$ that remains near the sea surface regardless of the island's vertical shift, reflecting coral growth in biologically viable depth ranges (typically 0-30 meters below sea level).

$$\tilde{\Phi}(\mathbf{p}) = (1 - \rho(\tilde{x}_{\mathbf{p}})) \cdot \Phi(\mathbf{p}). \quad (4)$$

This warp is applied to both the height field and the label map, ensuring consistent semantic deformation. For example, applying strokes to one side of a circular island creates concave coastlines while leaving high-resistance regions (e.g., the abyss) unaffected, simulating the localized impact of natural erosion (Figure 7, bottom).

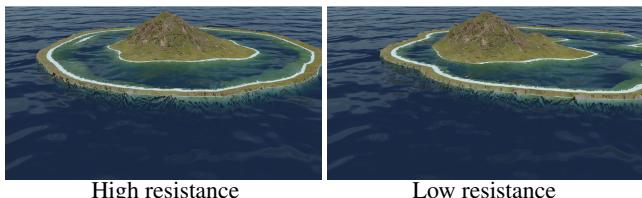


Figure 8: (Left) An island defined with a high resistance, under a uniform lateral wind velocity field, and (right) the same island with lower resistance on the reef borders, simulating the effect of coastal erosion.

4.2. Coral reef modeling

Once the terrain has been generated and deformed by the wind, we simulate the long-term geological evolution of coral reef islands through two parallel processes: the subsidence of the volcanic island and the upward growth of coral reefs. As observed in nature, the volcanic base sinks over time while coral formations grow vertically to remain close to the water surface, following the "keep-up" strategy of reef development.

4.2.1. Subsidence

Subsidence is modeled by uniformly scaling the original terrain height downward, simulating the gradual sinking of the volcanic

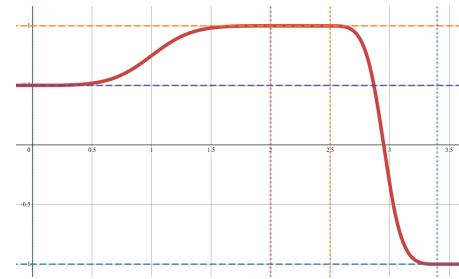


Figure 9: Our model describes reef growth by a piecewise function h_{coral} which is flat in the lagoon, the crest and abyss, and follows a smoothstep function as transitions for the backreef and fore reef regions. Zones' anchor heights are represented by horizontal dashed lines; zones' limits are dotted vertical lines.

We define distinct reef zones anchored at specific depths:

- Reef crest near sea level: $h_{\text{crest}} = -2 \text{ m}$,
- Back reef and lagoon: $h_{\text{back}} = -20 \text{ m}$,
- Fore reef sloping to abyss: $h_{\text{abyss}} = -100 \text{ m}$.

Each reef subregion is defined over a parametric domain $x \in [0, 1]$, with $x = 0$ the beginning of the reef region and $x = 1$ its end, directly inputting the parametric distance $x = \tilde{x} - i_{\text{reef}}$. For instance:

- Back reef: $x_{\text{back},\text{start}} = 0, x_{\text{back},\text{end}} = 0.5$,
- Reef crest: $x_{\text{crest},\text{start}} = 0.75, x_{\text{crest},\text{end}} = 0.8$,
- Abyss begins at $x_{\text{abyss},\text{start}} = 1$.

We model transition zones between these regions using a smoothstep operator:

$$\text{smoothstep}(x) = 3x^2 - 2x^3. \quad (6)$$

We denote the interpolating function as:

$$S(a, b, x_0, x_1, x) = a + (b - a) \text{smoothstep}\left(\frac{x - x_0}{x_1 - x_0}\right). \quad (7)$$

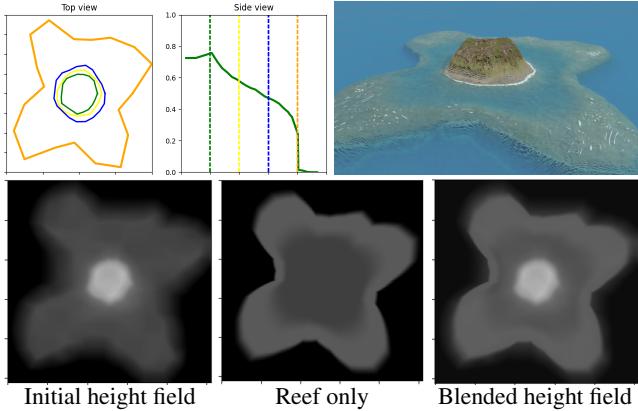


Figure 10: Volcano with single vent. Bottom left: the initial height field is computed directly from the user input. Bottom center: the reef height field is output from Equation (8). Bottom right: we blend the two results with Equation (9)."

301 The complete coral height field, as displayed in Figure 9, is built
302 as a piecewise function:

$$h_{\text{coral}}(x) = \sum_{r \in \text{subregions}} \begin{cases} h_r & \text{if } x_{r,\text{start}} \leq x \leq x_{r,\text{end}} \\ 0 & \text{otherwise} \end{cases} + \sum_{t \in \text{transitions}} \begin{cases} S(h_t, h_{t+1}, x_{t,\text{end}}, x_{t+1,\text{start}}, x) & \text{if } x_{t,\text{end}} < x < x_{t+1,\text{start}} \\ 0 & \text{otherwise} \end{cases} \quad (8)$$

4.2.3. Output

304 Finally, we merge the island base height field and the coral reef
305 height field by our *ad-hoc* smooth maximum operator smax defined
306 as:

$$\text{smax}(a, b) = \begin{cases} a + \frac{\delta}{2k} \left(\frac{1}{1+e^{-k\delta}} + \frac{1}{1-e^{-k\delta}} \right) & \text{for } a \neq b \\ a + \frac{1}{2k} & \text{for } a = b \end{cases} \quad (9)$$

307 with $\delta = b - a$ for conciseness, and k a sharpness parameter ap-
308 proximating the max operator as k increases. At $k = 5$, the operator
309 max is already well approximated while conserving continuity in
310 the resulting height field $\mathcal{H}(p) = \text{smax}(h_{\text{subsidi}}(\mathbf{p}), h_{\text{coral}}(\mathbf{p}))$.

311 The resulting terrain represents a plausible coral reef island,
312 where the volcanic island has subsided, and coral reefs have grown
313 upward to keep pace with the water level (Figure 10). The smooth
314 blending between the subsided terrain and the coral features en-
315 sures a natural transition between regions like the island, lagoon,
316 and coral reefs.

5. Data generation

317 The dataset is created using the procedural curve-based modeling
318 algorithm, with altered input parameters. For each generation, the
319 top-view and profile-view sketches use an initial layout. Each out-
320 line of the top-view sketch is defined as a centered circle of random
321

322 radius $r_{\min} \leq r^* \leq r_{\max}$. We add another deformation based on fBm
323 noise η such that the final contour, profile, and resistances, are de-
324 fined as

$$\begin{aligned} r(\theta) &= r^* + \eta_{\text{contours}}(\theta) \\ h_{\text{profile}}(\tilde{x}) &= h_{\text{profile}}^*(\tilde{x}) \cdot \eta_{\text{profile}}(\tilde{x}) \\ \rho(\tilde{x}) &= \rho^*(\tilde{x}) \cdot \eta_{\text{resistance}}(\tilde{x}) \end{aligned}$$

325 Finally, a wind velocity field is randomly generated as a final in-
326 put for the curve-based modeling algorithm, introducing deforma-
327 tions on the resulting islands. The realistic nature of wind is igno-
328 red for the generation of the wind strokes in order to provide com-
329 plexity and variety in the results. We generate a random number n of
330 strokes and their path by uniformly sampling a random number m
331 of points. The spread and intensity of each stroke is also random.

332 Once all inputs are set, we generate an example for multiple lev-
333 els of subsidence $\lambda \in [0, 1]$ to obtain a height field incorporating the
334 coral reef modeling and the associated label map.

335 We use the Hue component to encode the labels directly from
336 the parametric distance $H = \lfloor \tilde{x} \rfloor$ and encode the subsidence rate
337 into the Value component $V = \lambda$. Moreover, we purposefully left
338 the Saturation component unchanged at this stage, reserving space
339 for potentially including another parameter in the future.

Data augmentation

340 To enhance the variety of the dataset and improve the model's
341 ability to generalize, we apply several data augmentation tech-
342 niques in addition to the usual affine transformations (rotation, scal-
343 ing, and flipping):

344 *Translation:* Since the original algorithm always centers the is-
345 land, we translate the islands within the image to remove this con-
346 straint (Figure 11, leftmost). This ensures that the cGAN can gen-
347 erate islands in any position within the frame. By wrapping the island
348 on the borders, the model learns that data can appear on the edges
349 on the image.

350 *Copy-paste:* In some cases, we combine multiple islands into a
351 single sample, with only a maximum intersection over union (IoU)
352 of 10%, allowing the abysses to overlap but without obtaining other
353 region types to merge. Height fields blending is done through the
354 smooth maximum function from Equation (9), and the label blend-
355 ing uses the usual max operator. The regions not covered by any
356 island are assigned the abyss ID, which is very close to the min-
357 imal height, meaning that in this case specifically, the subsidence
358 factor (Value channel) is close to irrelevant (Figure 11, rightmost).

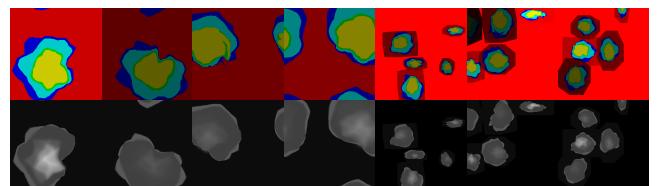


Figure 11: Examples from a dataset with increasing data augmen-
359 tation.

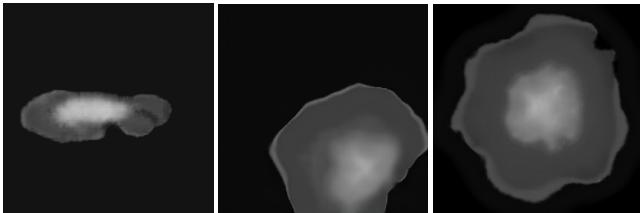


Figure 12: After the first epoch (left), grid artifacts similar to a low-resolution image are visible, but are being corrected during a second epoch (middle) where erosion patterns appear in the height fields (right).

360 All augmentation techniques are applied both to the height field
 361 and the label map simultaneously to ensure consistency between
 362 the input (the label map) and the output (the height field).

363 6. Learning-based generation

364 In this section, we introduce the use of a conditional Generative
 365 Adversarial Network (cGAN), specifically the pix2pix model pro-
 366 posed by Isola et al., 2017, to enhance the island generation process
 367 by increasing terrain variety and flexibility. While the curve-based
 368 modeling algorithm can create numerous island examples, cGAN
 369 provides additional flexibility in generating more complex terrain
 370 without the rigid constraints of the curve-based algorithm that stem
 371 from our initial assumptions based on coral reef formation theory.

372 The training was performed using a batch size of 1, and opti-
 373 mized using the Adam optimizer with a learning rate of 0.0002 and
 374 momentum parameters $\beta_1 = 0.5$ and $\beta_2 = 0.999$. The loss function
 375 combined a conditional adversarial loss and an L1 loss, where the
 376 L1 loss was weighted by a factor $\lambda = 100$. The generator follows
 377 a U-Net architecture with encoder-decoder layers and skip connec-
 378 tions, while the discriminator was based on a PatchGAN, classi-
 379 fying local image patches. These parameters are directly adopted
 380 from the original paper.

381 During inference, our model, with a dataset of approximatively
 382 10 000 examples representing about 30 minutes of training, still
 383 outputs grid artifacts, visible in Figure 12. After the second epoch,
 384 visual artifacts are already rare. This training time seems long but,
 385 at the best of our knowledge, is the shortest training time for terrain
 386 generation with a learning-based approach in the literature, mostly
 387 due to the synthetic nature of our dataset.

388 Once the model is trained, the user colors a 256×256 RGB im-
 389 age to sketch the different regions. Each region is given a specific
 390 color based on the HSV encoding used for the dataset generation.
 391 The examples presented in this paper use red for abysses, blue for
 392 reefs, cyan for lagoons, green for beaches and yellow for moun-
 393 tains. The subsidence factor presented in Section 4.2.1 is control-
 394 lable through the luminosity of the input image.

395 The Python script for the initial island dataset generation is unop-
 396 timized and takes about 2.5s per island of size 256×256 as we do
 397 not perform parallelization here. Implementing an optimized C++

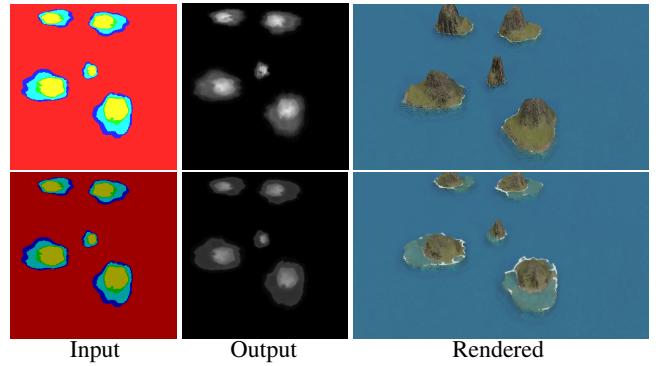


Figure 13: An identical label map yields similar height fields over multiple inferences from the model, even after modifying the subsidence factor (visible in the luminosity of the input image).

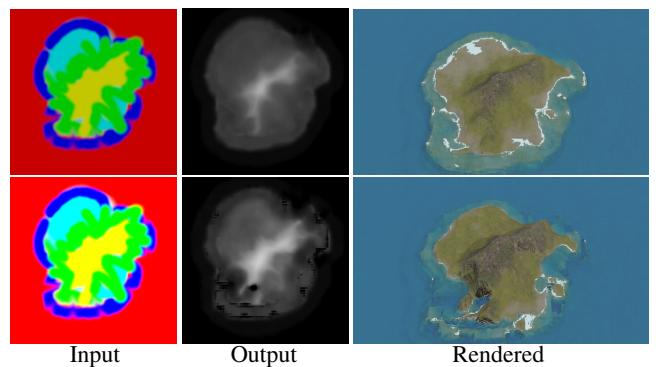


Figure 14: User applied fuzzy brushes to draw the label map, resulting in some pixels that are inconsistent with the dataset and unlogical island layouts: abyss regions (red) are found between beach (green), lagoon (cyan) and reef regions (blue). The neural model ignores the layout inconsistencies, and partially over-brightness as long as the pixel isn't completely white.

398 version of the initial generation process reduces this execution time
 399 to 50ms per generation.

400 On the other hand, the inference time of our deep learning model
 401 for a single input image of dimension 256×256 remains constant
 402 regardless of scene complexity. Using the NVIDIA GeForce GTX
 403 1650 Ti GPU with Python 3.10 and PyTorch version 2.5.1+cu121,
 404 the inference time measured is 5ms (std 1.1ms).

405 7. Results

406 The resulting model for coral island generation enables a high
 407 control-level from a user perspective as the unconstrained painting
 408 allows for complex scenarios while producing in real-time the
 409 resulting height fields. In this paper we used the software Blender
 410 to provide renders directly from the output height fields. As our
 411 pix2pix model is trained to output 256×256 images, the resolution
 412 of the 3D models is limited by this architecture.

413 Using deep-learning-based models, most constraints from our
 414 initial assumptions are lifted (radial layout, isolated islands, ...).

415 The control over the overall shapes of the islands regions are given
 416 through digital painting, here using the GIMP software. Each pixel
 417 of the image is encoded in HSV, with the region identifier encoded
 418 in the Hue channel. The user may increase or decrease the subsi-
 419 dence level of the island by modifying the Value channel over the
 420 whole image (see Figure 13).

421 Since the model relies on pixel-level statistics rather than strict
 422 class values, it is robust to noise caused by compression, anti-
 423 aliasing from brush tools, or resizing artifacts introduced by im-
 424 age editors. The example displayed in Figure 14 presents a sketch
 425 for which the outlines of the regions are at the same time blurry
 426 and with layouts that are not expected (such as the small red re-
 427 gions inside the southern lagoon region or the adjacency of beach
 428 regions directly with the abyssal region) on the top figure and show-
 429 ing highlight clipping on the bottom figure. The learned model does
 430 not include inconsistencies and results in plausible 3D models. We
 431 can note that the input label map isn't require to be hand-drawn,
 432 as a simple Perlin noise can produce it randomly, as shown in the
 433 examples of Figure 17.

434 The tolerance over the input values may be used to provide even
 435 more control about the transitions between two regions. Figure 15
 436 shows an example of input map with regions that are leaking over
 437 neighboring regions, and the introduction of new hue values nonex-
 438 istent in the dataset (light green and dark green) but are the inter-
 439 polated hue value of mountain regions and beach regions.

440 Since the curve-based procedural phase included low random-
 441 ness, the output of the cGAN is limiting its unpredictability, mean-
 442 ing that small input changes results only in minor changes on the
 443 output, preventing unexpected results. Figure 13 shows the result
 444 of an input map with only a variation on the subsidence level, the
 445 resulting height fields are very similar. Adding the real-time com-
 446 putation of outputs, it becomes possible to construct progressively
 447 a landscape and correct small mistakes to intuitively design islands
 448 inspired by real-world regions. A creation of an island similar to
 449 Mayotte, presented in Figure 16 was hand-made in just a few min-
 450 utes.

451 **Limitations** While this approach brings significant advantages,
 452 there are also some limitations to consider. The reliance on a syn-
 453 thetic dataset means that the cGAN inherits some biases and limita-
 454 tions of our initial curve-based algorithm. This could limit the true
 455 diversity of the terrains that the model can generate, as the output
 456 is confined by the patterns present in the training data. Addition-

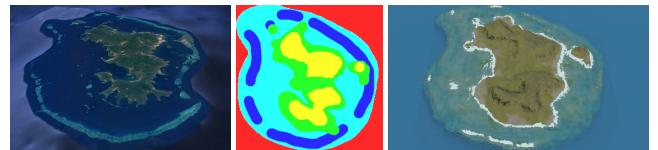


Figure 16: Comparison between real (left, from Google Earth) and synthetic (right) islands of Mayotte. The label map (center) is hand-drawn to match approximatively its real-world counterpart.

457 ally, the cGAN model's internal logic lacks transparency, offering
 458 limited user control over the generation process once the model has
 459 been trained. Moreover, the training time is far from real-time.

460 **Future work** Further improvements could be made to the synthetic
 461 dataset. Incorporating more complex geological processes, such as
 462 wave erosion or tidal influences, could lead to even more realistic
 463 terrains. Additionally, refining the way islands are blended in multi-
 464 island samples, or adding more diverse input conditions (e.g., dif-
 465 ferent geological settings), could help the model generalize better
 466 and produce more varied and dynamic landscapes. While the cur-
 467 rent model allows for rapid terrain generation, adding more options
 468 for users to interact with the cGAN, such as tweaking parameters
 469 like wind strength or island size, could enhance the flexibility of
 470 the system. Many other neural network models could be exploited
 471 to increase the possibilities, such as newer variants of cGANs (Park
 472 et al., 2019), or models with style transfer functionalities (Gatys et
 473 al., 2015; Zhu et al., 2020) in order to change the overall aspect
 474 of a terrain (Perche et al., 2023b, 2023a), use text-to-images mod-
 475 els (Rombach et al., 2021; Radford et al., 2021) to generate height
 476 fields from a verbal prompt, or super-resolution models (Dong et
 477 al., 2014) to increase the definition of details in the final output
 478 (Guérin, Digne, Galin, and Peytavie, 2016).

479 8. Conclusion

480 We presented a novel approach to generating coral reef island ter-
 481 rains by combining traditional procedural methods with deep learn-
 482 ing techniques. We first developed a procedural generation algo-
 483 rithm capable of creating a wide variety of island terrains using
 484 top-view and profile-view sketches, wind deformation, subsidence,
 485 and coral reef growth simulation. By applying these methods, we
 486 were able to produce realistic terrains based on geological pro-
 487 cesses, capturing key features of coral reef islands such as beaches,
 488 lagoons, and reefs.

489 To further enhance flexibility and realism in the generation pro-
 490 cess, we incorporated a conditional Generative Adversarial Net-
 491 work (cGAN), using the pix2pix model to translate island feature
 492 label maps into height maps. The cGAN model helped to overcome
 493 some of the constraints inherent in the procedural algorithm, such
 494 as radial symmetry and fixed island positioning. Using data aug-
 495 mentation techniques, we trained the cGAN on a synthetic dataset
 496 to produce varied and realistic island terrains.

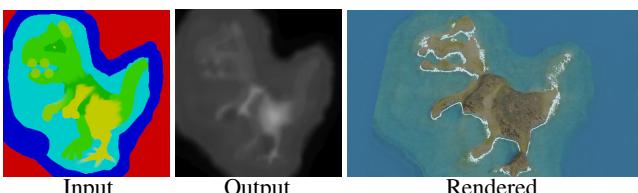


Figure 15: Without constraints on the generation, the user may use unrealistic layout and the neural network will however output a plausible result. Here new colors have been added (pistachio), but the network naturally considers it as a transition from mountain to beach.

- 497 **References**
- 498 Beckham, C., & Pal, C. (2017). A step towards procedural terrain
499 generation with gans. <http://arxiv.org/abs/1707.03383>
500 (cit. on p. 2).
- 501 Beneš, B., Těšínský, V., Hornyš, J., & Bhatia, S. K. (2006). Hy-
502 draulic erosion. *Computer Animation and Virtual Worlds*,
503 17, 99–108. <https://doi.org/10.1002/cav.77> (cit. on p. 2).
- 504 Cordonnier, G., Braun, J., Cani, M.-P., Benes, B., Galin, É., Pey-
505 tavie, A., & Guérin, É. (2016). Large scale terrain gener-
506 ation from tectonic uplift and fluvial erosion. *Computer*
507 *Graphics Forum*, 35, 165–175. <https://doi.org/10.1111/cgf.12820> (cit. on p. 2).
- 509 Cordonnier, G., Cani, M.-P., Beneš, B., Braun, J., & Galin, É.
510 (2017). Sculpting mountains: Interactive terrain model-
511 ing based on subsurface geology. *IEEE Transactions on*
512 *Visualization and Computer Graphics*, 24. <https://doi.org/10.1109/TVCG.2017.2689022> (cit. on p. 2).
- 514 Cordonnier, G., Galin, É., Gain, J., Beneš, B., Guérin, É., Peytavie,
515 A., & Cani, M.-P. (2017). Authoring landscapes by com-
516 bining ecosystem and terrain erosion simulation. *ACM*
517 *Transactions on Graphics*, 36. <https://doi.org/10.1145/3072959.3073667> (cit. on p. 2).
- 519 Dong, C., Loy, C. C., He, K., & Tang, X. (2014). Image super-
520 resolution using deep convolutional networks. <http://arxiv.org/abs/1501.00092> (cit. on p. 8).
- 522 Ebert, D. S., Peachey, D., Worley, S., Hart, J. C., Musgrave, K.,
523 Perlin, K., & Mark, W. R. (2003). *Texturing and model-
524 ing, a procedural approach* (Vol. Third edition). Morgan
525 Kaufmann. (Cit. on p. 2).
- 526 Ecormier-Nocca, P., Cordonnier, G., Carrez, P., Moigne, A. M.,
527 Memari, P., Benes, B., & Cani, M. P. (2021). Authoring
528 consistent landscapes with flora and fauna. *ACM Trans-
529 actions on Graphics*, 40. <https://doi.org/10.1145/3450626.3459952> (cit. on p. 2).
- 531 Fournier, A., Fussell, D., & Carpenter, L. (1982). Computer ren-
532 dering of stochastic models. *Communications of the ACM*,
533 25, 371–384. <https://doi.org/10.1145/358523.358553>
534 (cit. on p. 2).
- 535 Gain, J., Marais, P., & Straßer, W. (2009). Terrain sketching. *Pro-
536 ceedings of I3D 2009: The 2009 ACM SIGGRAPH Sym-
537 posium on Interactive 3D Graphics and Games*, 1, 31–
538 38. <https://doi.org/10.1145/1507149.1507155> (cit. on
539 p. 2).
- 540 Galin, E., Guérin, E., Peytavie, A., Cordonnier, G., Cani, M.-P.,
541 Benes, B., & Gain, J. (2019). A review of digital ter-
542 rain modeling. *Computer Graphics Forum*, 38, 553–577.
543 <https://doi.org/10.1111/cgf.13657> (cit. on p. 2).
- 544 Gasch, C., Chover, M., Remolar, I., & Rebollo, C. (2020). Pro-
545 cedural modelling of terrains with constraints. *Multimedia*
546 *Tools and Applications*, 79, 31125–31146. <https://doi.org/10.1007/s11042-020-09476-3> (cit. on p. 2).
- 548 Gatys, L. A., Ecker, A. S., & Bethge, M. (2015). A neural algorithm
549 of artistic style. <http://arxiv.org/abs/1508.06576> (cit. on
550 p. 8).
- 551 Génevaux, J.-D., Galin, É., Peytavie, A., Guérin, É., Briquet, C.,
552 Grosbellet, F., & Beneš, B. (2015). Terrain modelling
553 from feature primitives. <https://doi.org/https://doi.org/10.1111/cgf.12530> (cit. on p. 2).
- 555 Guérin, E., Peytavie, A., Masnou, S., Digne, J., Sauvage, B., Gain,
556 J., & Galin, E. (2022). Gradient terrain authoring. *Com-
557 puter Graphics Forum*, 41, 85–95. <https://doi.org/10.1111/cgf.14460> (cit. on p. 2).
- 558 Guérin, É., Digne, J., Galin, É., & Peytavie, A. (2016). Sparse
559 representation of terrains for procedural modeling. *Com-
560 puter Graphics Forum*, 35, 177–187. <https://doi.org/10.1111/cgf.12821> (cit. on p. 8).
- 561 Guérin, É., Digne, J., Galin, É., Peytavie, A., Wolf, C., Beneš,
562 B., & Martinez, B. (2017). Interactive example-based
563 terrain authoring with conditional generative adversar-
564 ial networks. *ACM Transactions on Graphics*, 36. <https://doi.org/10.1145/3130800.3130804> (cit. on p. 2).
- 565 Hnaidi, H., Guérin, E., Akkouche, S., Peytavie, A., & Galin, E.
566 (2010). Feature based terrain generation using diffusion
567 equation. *Computer Graphics Forum*, 29, 2179–2186.
568 <https://doi.org/10.1111/j.1467-8659.2010.01806.x>
569 (cit. on p. 2).
- 570 Hopley, D. (2014). *Encyclopedia of modern coral reefs : Structure,
571 form and process*. Springer, Credo Reference. (Cit. on
572 p. 1).
- 573 Isola, P., Zhu, J.-Y., Zhou, T., & Efros, A. A. (2017). Image-to-
574 image translation with conditional adversarial networks.
575 *2017 IEEE Conference on Computer Vision and Pattern*
576 *Recognition (CVPR)*, 5967–5976. <https://doi.org/10.1109/CVPR.2017.632> (cit. on pp. 2, 7).
- 577 Körner, C., Spehn, E. M., Bugmann, H., & Zurich, E. (2014).
578 *Mountain systems* (tech. rep.). <https://www.researchgate.net/publication/238663492> (cit. on p. 1).
- 579 Lyons, M. B., Murray, N. J., Kennedy, E. V., Kovacs, E. M.,
580 Castro-Sanguino, C., Phinn, S. R., Acevedo, R. B., Al-
581 varez, A. O., Say, C., Tudman, P., Markey, K., Roe, M.,
582 Canto, R. F., Fox, H. E., Bambic, B., Lieb, Z., Asner,
583 G. P., Martin, P. M., Knapp, D. E., ... Roelfsema, C. M.
584 (2024). New global area estimates for coral reefs from
585 high-resolution mapping. *Cell Reports Sustainability*, 1,
586 100015. <https://doi.org/10.1016/j.crsus.2024.100015>
587 (cit. on p. 1).
- 588 Mei, X., Decaudin, P., & Hu, B. G. (2007). Fast hydraulic ero-
589 sion simulation and visualization on gpu. *Proceedings -
590 Pacific Conference on Computer Graphics and Applica-
591 tions*, 47–56. <https://doi.org/10.1109/PG.2007.27> (cit. on
592 p. 2).
- 593 Mirza, M., & Osindero, S. (2014). Conditional generative adversar-
594 ial nets. <http://arxiv.org/abs/1411.1784> (cit. on p. 2).
- 595 Musgrave, F. K., Kolb, C. E., & Mace, R. S. (1989). The synthesis
596 and rendering of eroded fractal terrains. *Proceedings of*
597 *the 16th Annual Conference on Computer Graphics and*
598 *Interactive Techniques, SIGGRAPH 1989*, 41–50. <https://doi.org/10.1145/74333.74337> (cit. on p. 2).
- 599 Natali, M., Viola, I., & Patel, D. (2012). Rapid visualization of
600 geological concepts. *Brazilian Symposium of Computer*
601 *Graphic and Image Processing*, 150–157. <https://doi.org/10.1109/SIBGRAPI.2012.29> (cit. on p. 2).
- 602 Neidhold, B., Wacker, M., & Deussen, O. (2005). Interactive phys-
603 ically based fluid and erosion simulation. *Natural Phe-*
604 *nomena*, 1, 1–10.

- 611 *nomena*, 25–32. <http://graphics.uni-konstanz.de/~publikationen/Neidhold2005InteractivePhysicallyBased/Neidhold2005InteractivePhysicallyBased.pdf> (cit. on p. 2).
- 612 Olsen, J. (2004). Realtime procedural terrain generation. *Department of Mathematics And Computer Science*, 20. <https://pdfs.semanticscholar.org/5961/c577478f21707dad53905362e0ec4e6ec644.pdf> (cit. on p. 2).
- 613 Panagiotou, E., & Charou, E. (2020). Procedural 3d terrain generation using generative adversarial networks. <https://arxiv.org/abs/2010.06411> (cit. on p. 2).
- 614 Park, J., Redwine, J., Hill, T. D., & Kotun, K. (2019). Water resource and ecotone transformation in coastal ecosystems. *Ecological Modelling*, 405, 69–85. <https://doi.org/10.1016/j.ecolmodel.2019.04.015> (cit. on p. 8).
- 615 Perche, S., Peytavie, A., Benes, B., Galin, E., & Guérin, E. (2023a). Styledem: A versatile model for authoring terrains. <https://arxiv.org/abs/2304.09626> (cit. on p. 8).
- 616 Perche, S., Peytavie, A., Benes, B., Galin, E., & Guérin, E. (2023b). Authoring terrains with spatialised style. *Computer Graphics Forum*, 42. <https://doi.org/10.1111/cgf.14936> (cit. on p. 8).
- 617 Perlin, K. (1985). An image synthesizer. *ACM SIGGRAPH Computer Graphics*, 19, 287–296. <https://doi.org/10.1145/325165.325247> (cit. on p. 2).
- 618 Perlin, K. (2001). Noise hardware. *Real-Time Shading SIGGRAPH Course Notes* (cit. on p. 2).
- 619 Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, A., Mishkin, P., Clark, J., Krueger, G., & Sutskever, I. (2021). Learning transferable visual models from natural language supervision. <https://arxiv.org/abs/2103.00020> (cit. on p. 8).
- 620 Rombach, R., Blattmann, A., Lorenz, D., Esser, P., & Ommer, B. (2021). High-resolution image synthesis with latent diffusion models. <https://arxiv.org/abs/2112.10752> (cit. on p. 8).
- 621 Schott, H., Paris, A., Fournier, L., Guérin, E., & Galin, E. (2023). Large-scale terrain authoring through interactive erosion simulation. *ACM Transactions on Graphics*, 42, 1–15. <https://doi.org/10.1145/3592787> (cit. on p. 2).
- 622 Sisodia, Y. (2022). Gan-generated terrain for game assets. *Indian Journal of Artificial Intelligence and Neural Networking*, 2, 1–3. <https://doi.org/10.54105/ijainn.F1060.102622> (cit. on p. 2).
- 623 Smelik, R. M., Kraker, K. J. D., Groenewegen, S. A., Tutenel, T., & Bidarra, R. (2009). A survey of procedural methods for terrain modelling. *Proceedings of the CASA workshop on 3D advanced media in gaming and simulation (3AMIGAS)* (cit. on p. 2).
- 624 Spick, R., & Walker, J. A. (2019). Realistic and textured terrain generation using gans. *Proceedings - CVMP 2019: 16th ACM SIGGRAPH European Conference on Visual Media Production*. <https://doi.org/10.1145/3359998.3369407> (cit. on p. 2).
- 625 Talgorn, F. X., & Belhadj, F. (2018). Real-time sketch-based terrain generation. *ACM International Conference Proceeding Series*, 13–18. <https://doi.org/10.1145/3208159.3208184> (cit. on p. 2).
- 626 Voulgaris, G., Mademlis, I., & Pitas, I. (2021). Procedural terrain generation using generative adversarial networks. *European Signal Processing Conference, 2021-August*, 686–690. <https://doi.org/10.23919/EUSIPCO54536.2021.9616151> (cit. on p. 2).
- 627 Wulff-Jensen, A., Rant, N. N., Møller, T. N., & Billeskov, J. A. (2018, January). Deep convolutional generative adversarial network for procedural 3d landscape generation based on dem. In *Interactivity, game creation, design, learning, and innovation* (pp. 85–94). Springer. https://doi.org/10.1007/978-3-319-76908-0_9 (cit. on p. 2).
- 628 Zhu, D., Cheng, X., Zhang, F., Yao, X., Gao, Y., & Liu, Y. (2020). Spatial interpolation using conditional generative adversarial neural networks. *International Journal of Geographical Information Science*, 34, 735–758. <https://doi.org/10.1080/13658816.2019.1599122> (cit. on p. 8).

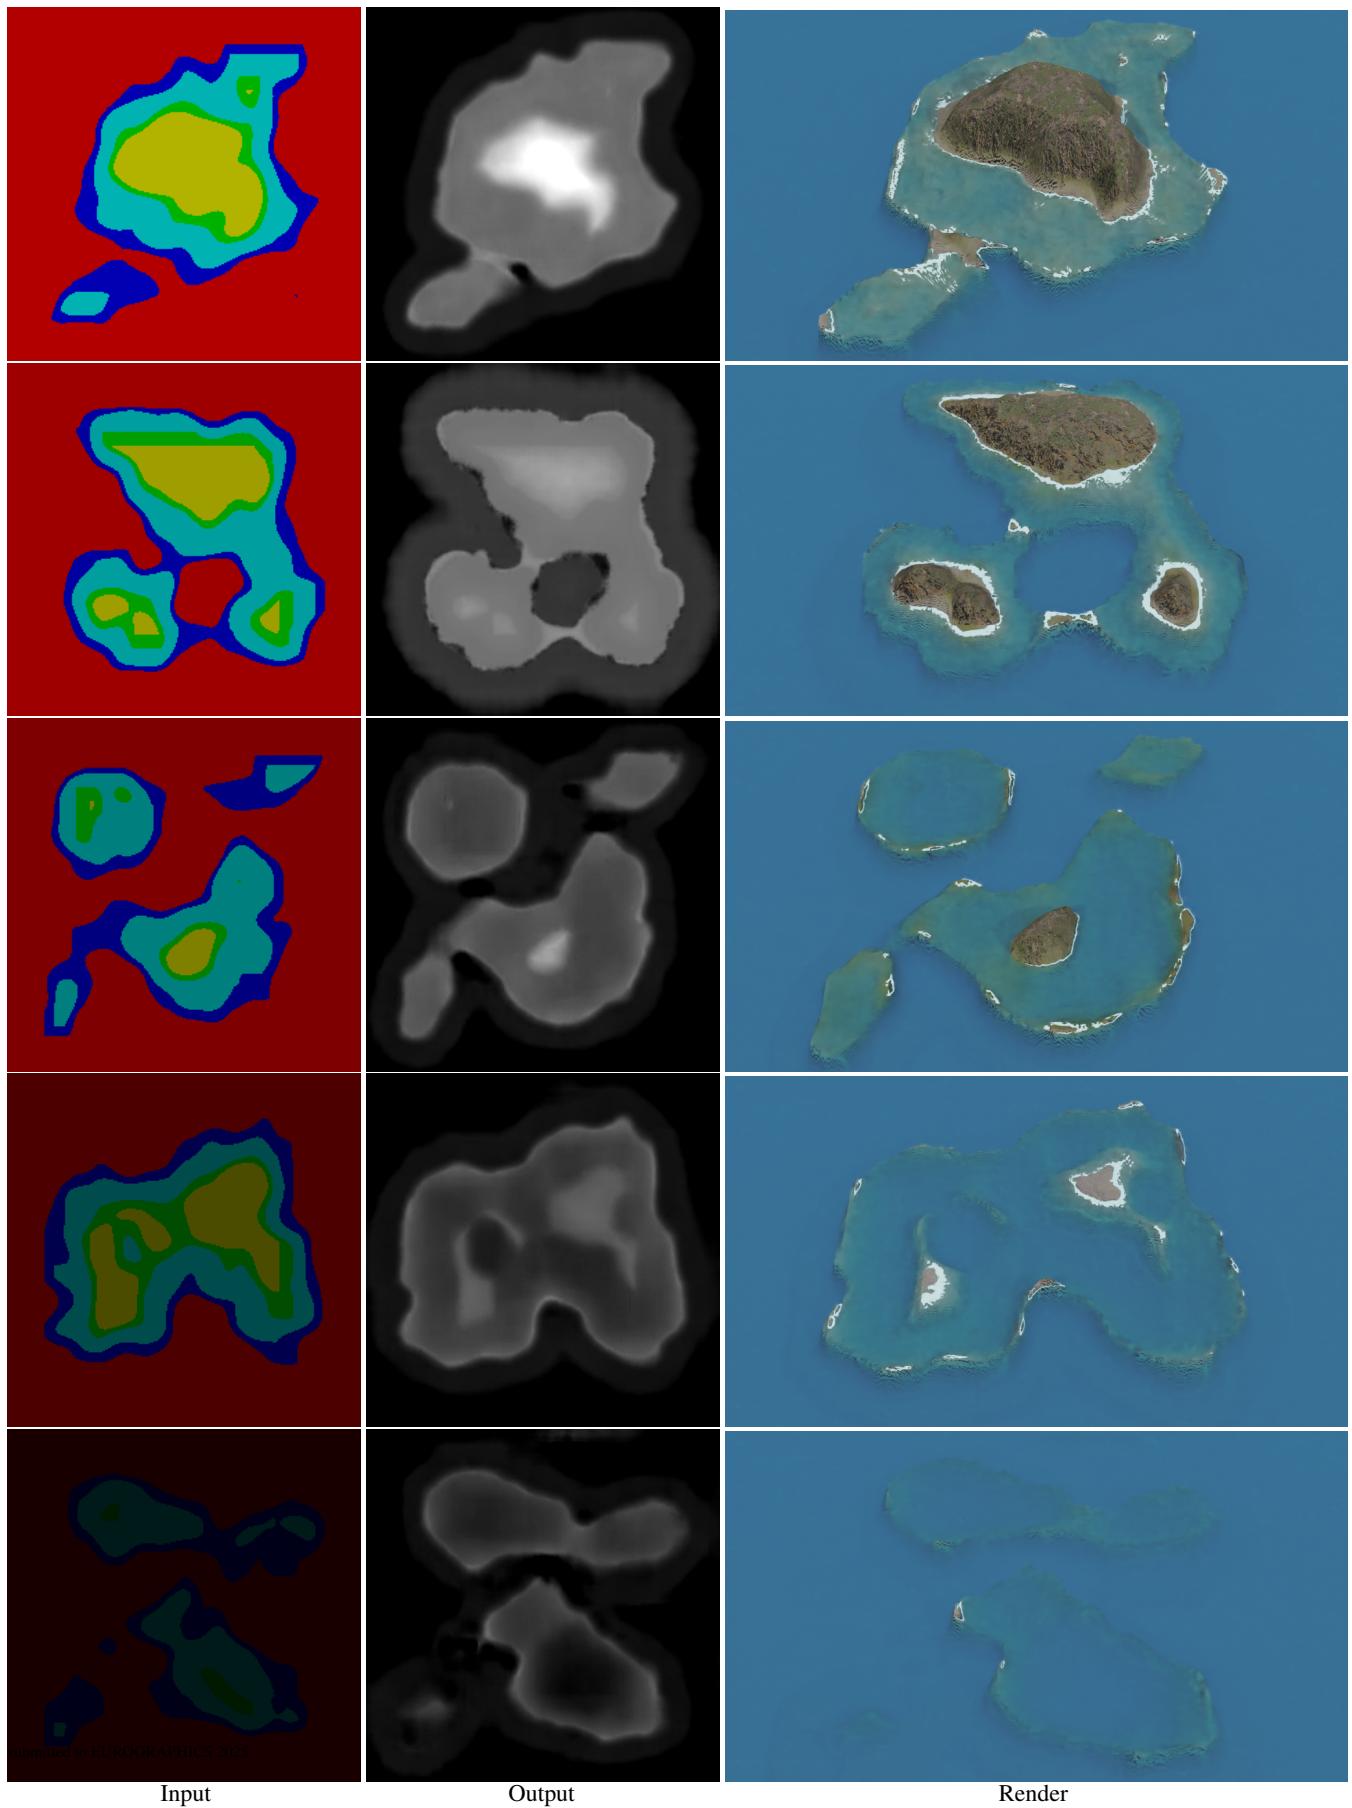


Figure 17: Starting from random Perlin noise, transformed into label maps with increasing subsidence, we can generate a large variety of results.

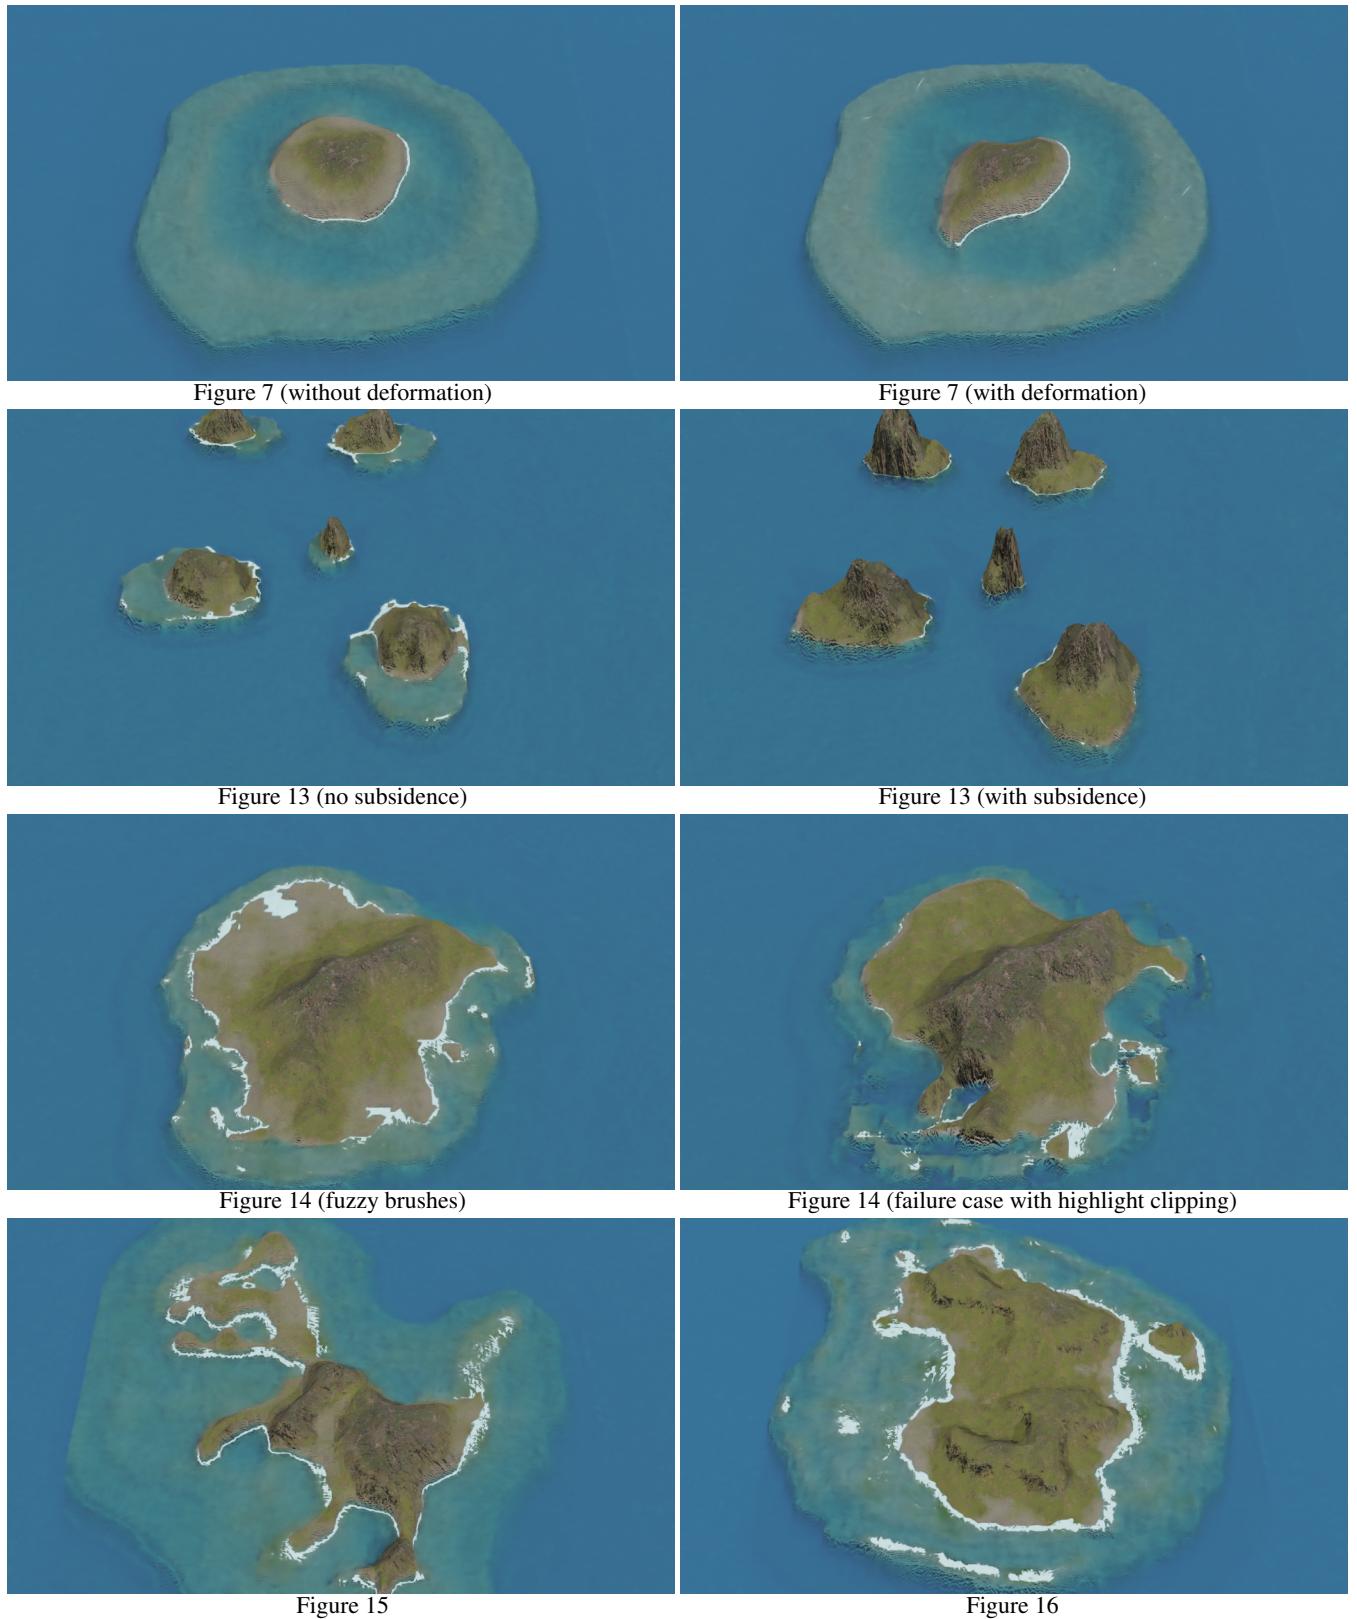


Figure 18: High-resolution version of figures found in this paper.